



## Module 01: Kafka Core




# Course Agenda

- ✦ Module 01 – Kafka Core
- ✦ Module 02 – Installing Kafka
- ✦ Module 03 – Kafka CLI
- ✦ Module 04 – Programmatic API
- ✦ Module 05 – Advanced Programming
- ✦ Module 06 – Kafka Streams Overview
- ✦ Module 07 – Kafka EchoSystem and Administration

# Agenda

- ✦ Kafka Brokers
- ✦ Topics and Partitions
- ✦ Producers
- ✦ Consumers and Consumer Groups

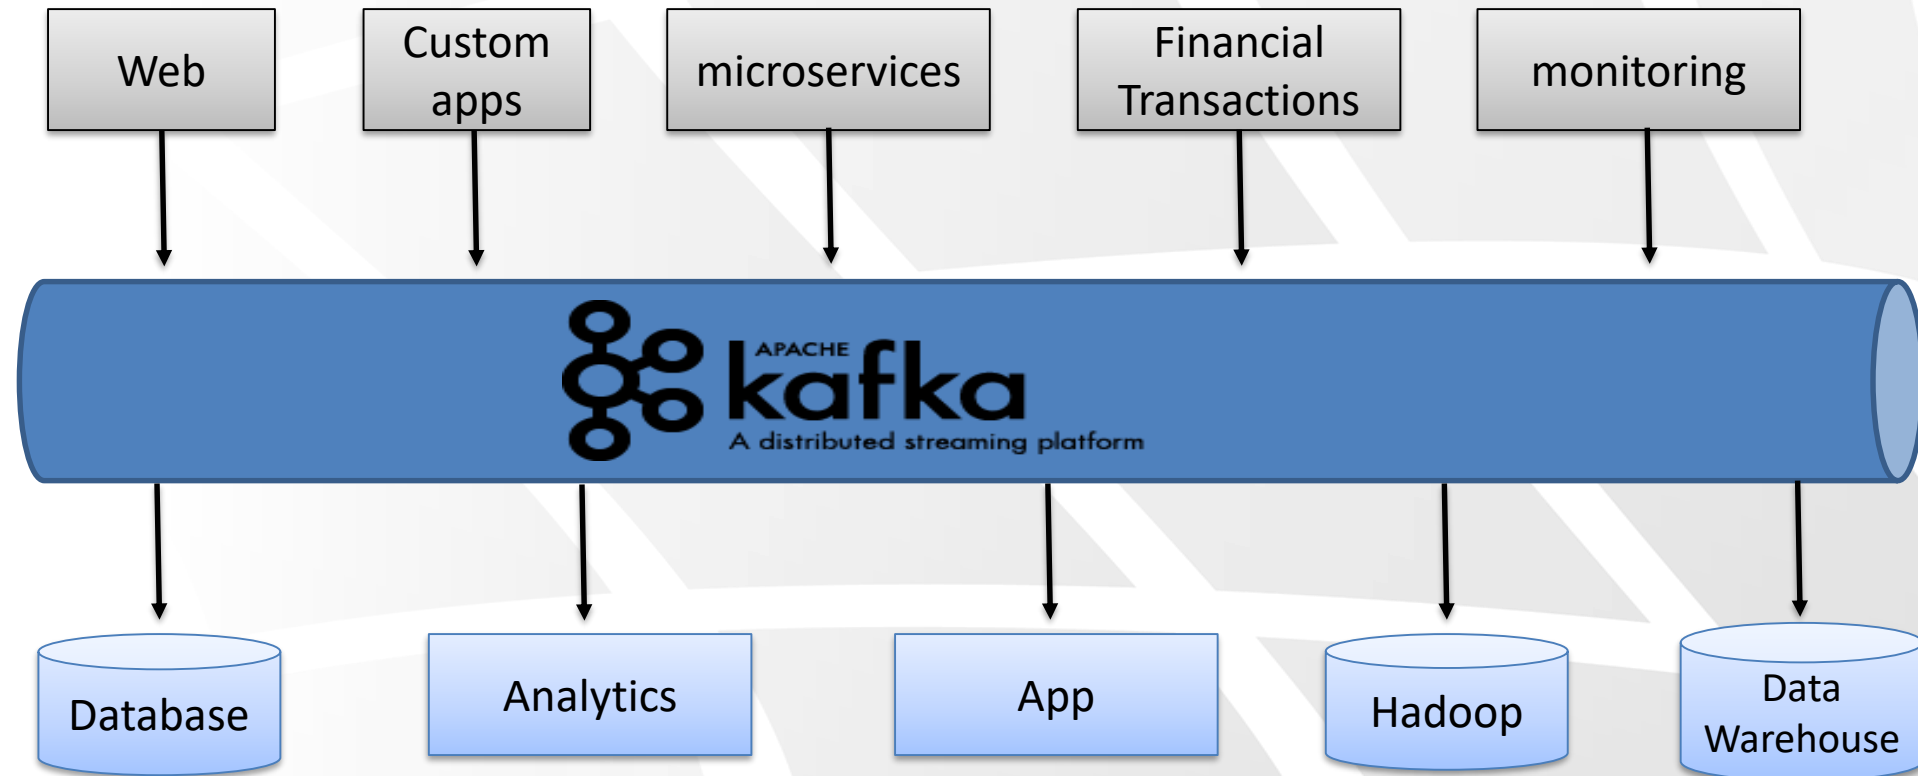
# Defining the Decoupling problem



- ✦ 4 sources + 4 targets mean 16 integrations
- ✦ Protocols (TCP/Http/Rest/FTP)
- ✦ Data Format – how the data is parsed (Binary, Json, Avro, CSV)
- ✦ Data Schema & evolution – how a data is shaped and may change

# Apache Kafka

✦ A high throughput distributed system



# Why Apache Kafka

- ✦ Created by LinkedIn, now open source promoted by confluence
- ✦ Distributed
- ✦ Fault tolerant
- ✦ Resilient Architecture
- ✦ Scales Horizontally :
  - ✦ Can scale to hundreds of brokers
  - ✦ Can scale to millions of messages per second
- ✦ High Performance
- ✦ Low latency (less than 10ms)
- ✦ Used by 2000+ firms ,35% of the Fortune 500



# Use Cases

- ✦ Messaging System – Kafka has better throuput, build-in partitions, replication and fault-tolerance than most message broker systems.
- ✦ Activity Tracking – rebuild user activity pipeline as a pub/sub feeds
- ✦ Metrics gathering – operational monitoring data for distributed apps.
- ✦ Log Aggregation - cleaner abstractions of log as stream of data.
- ✦ Stream Processing – data processing in multiple stages (Kafka streams)
- ✦ Event Sourcing – large store log enables time-order sequence

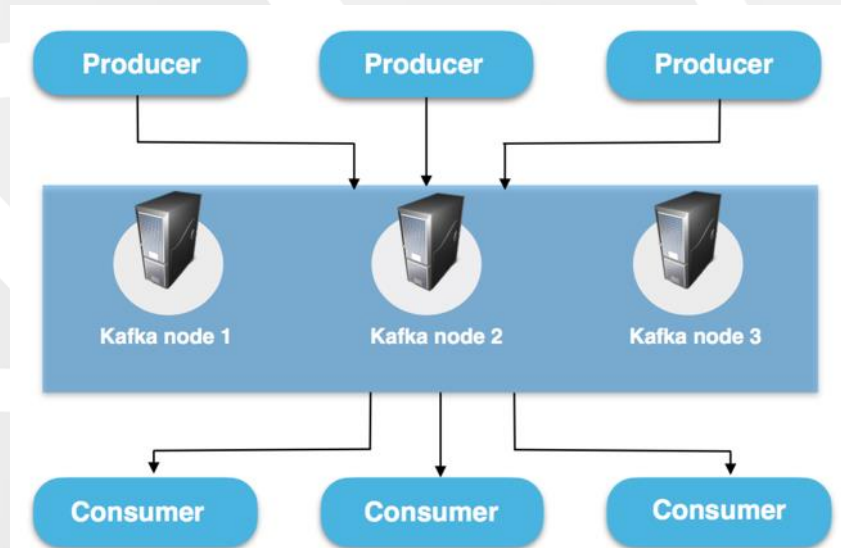
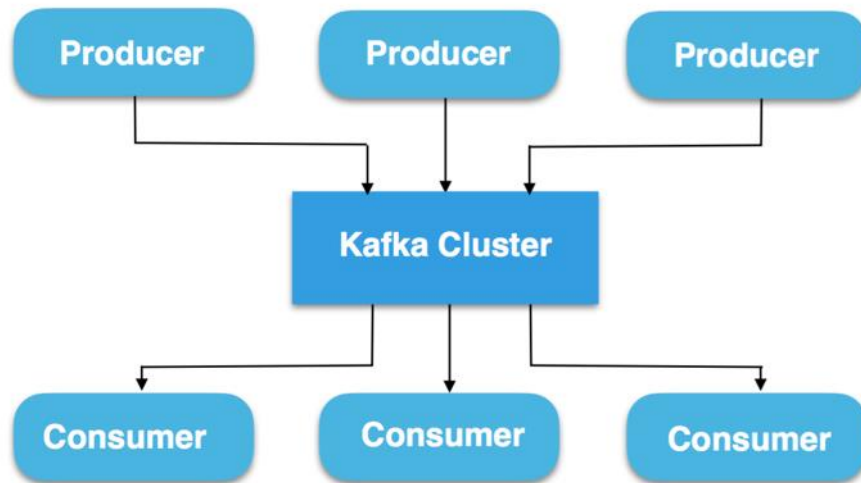
# Use Cases

- ✦ **NETFLIX** uses Kafka to apply recommendations in real time while you're watching TV shows
- ✦ **UBER** uses kafka to gather user, taxi and trip data in real time to capture and forecast demand and compute surge pricing
- ✦ **Linkedin** uses Kafka to prevent spam, collect user interactions to make better connection recommendation in real time.




# Brokers

- ✦ A Kafka cluster consists one or more servers (a.k.a 'brokers')
- ✦ **Producers** are processes that publish data (push messages) into Kafka topic within the broker
- ✦ **Consumers** are processes that pulls (read) data off a Kafka topic



# Topic


- ✦ A particular stream of Data
- ✦ Can be think of as a 'Category' or a table in a Database (without all the constraints)
- ✦ Each Topic has a unique name
- ✦ Producers and Consumer write/read from a specific topic



# Anatomy of a Topic

- ⚡ Topics Are Split into partitions :
- ⚡ Each Partition is ordered
- ⚡ Each message within a partition gets an incremental Id called offset
- ⚡ Data in partitions can be retained for a configurable amount of time (default = One Week)

## Anatomy of a Topic

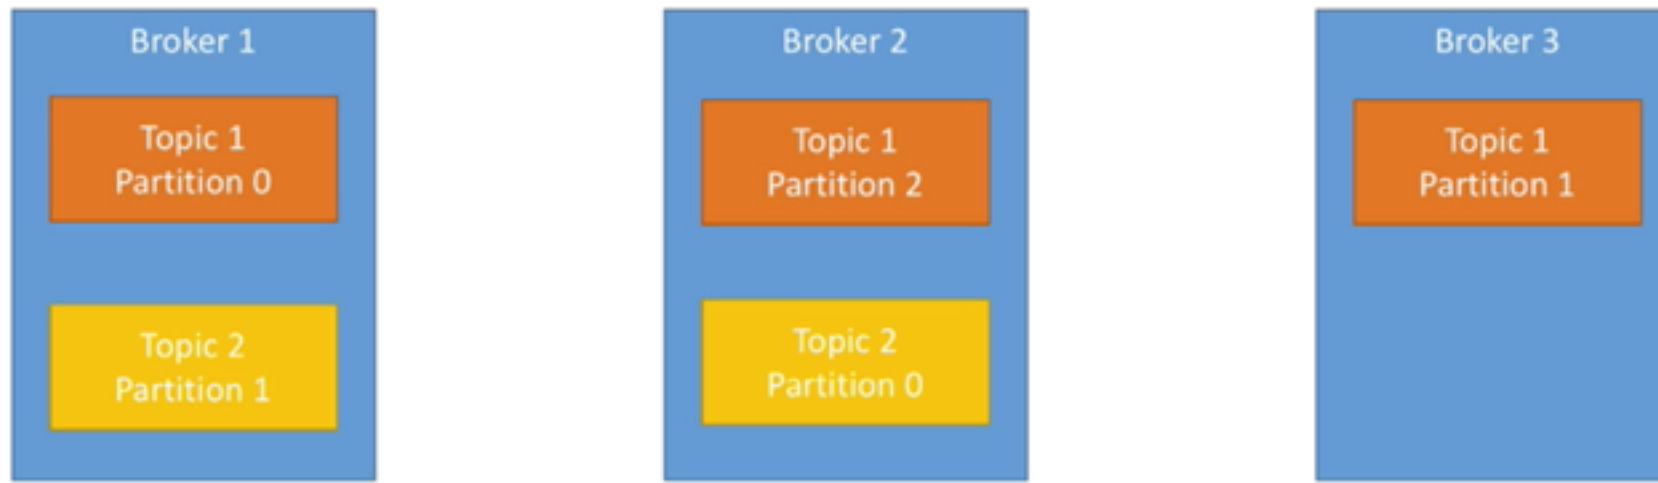


# Anatomy of a Topic

- ✦ Once data is written to a topic it cannot be changed (immutability)
- ✦ Data is assigned randomly to a partition unless a key is specified
- ✦ You can have as many partitions per topic as you want
- ✦ Order is guaranteed only within a partition


# Brokers

- ✦ A Kafka cluster is composed of multiple brokers (servers)
- ✦ Each broker contains certain topic partitions
- ✦ After connecting to any broker (called bootstrap server), you will be connect to the entire cluster




# Topic Replication Factor

- ✦ Replication factor determines how data is replicated across the nodes
- ✦ This allows Kafka to automatically failover to the replica when a server in the cluster fails
- ✦ Replica happens at the partition granularity
  - ✦ Example : Topic with 2 partitions and replication factor of 2




# Topic Replication Factor

✦ Example : Losing broker 2 can still serve the data



# Topic Leader

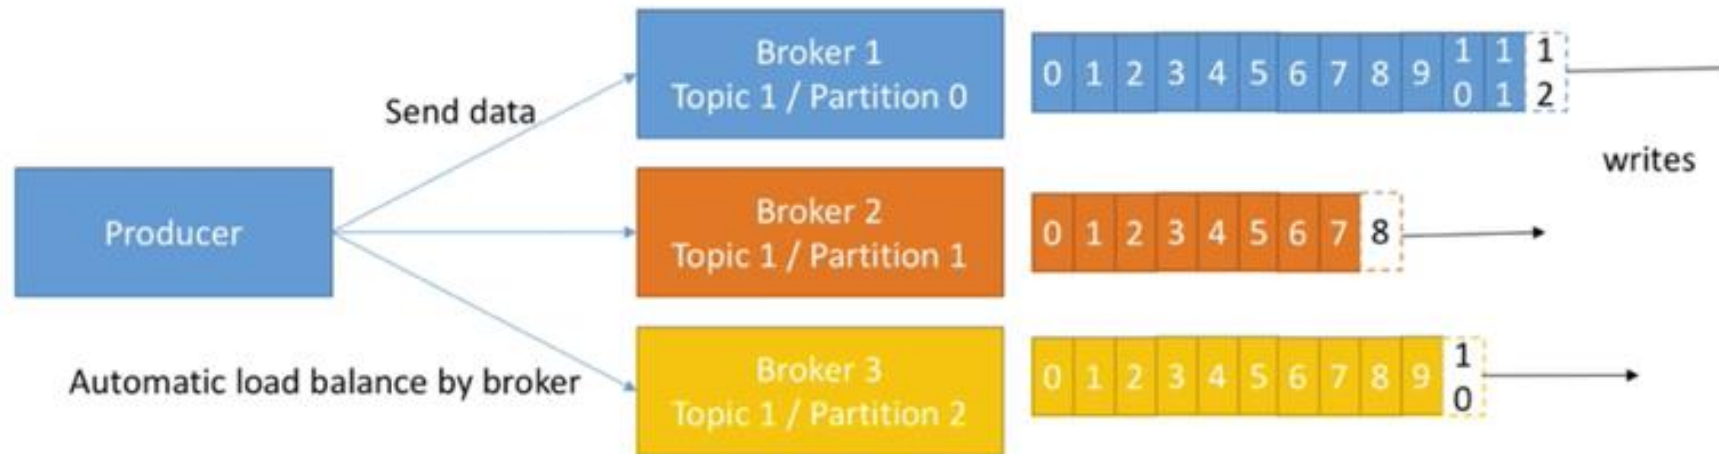
- ✦ One replica is designated as the “leader” while other are followers.
- ✦ At a given time only 1 broker can be a leader for a given partition
- ✦ Only that leader can send and receive data
- ✦ The other partitions will sync to the leader, also called **ISR** - In Sync Replica





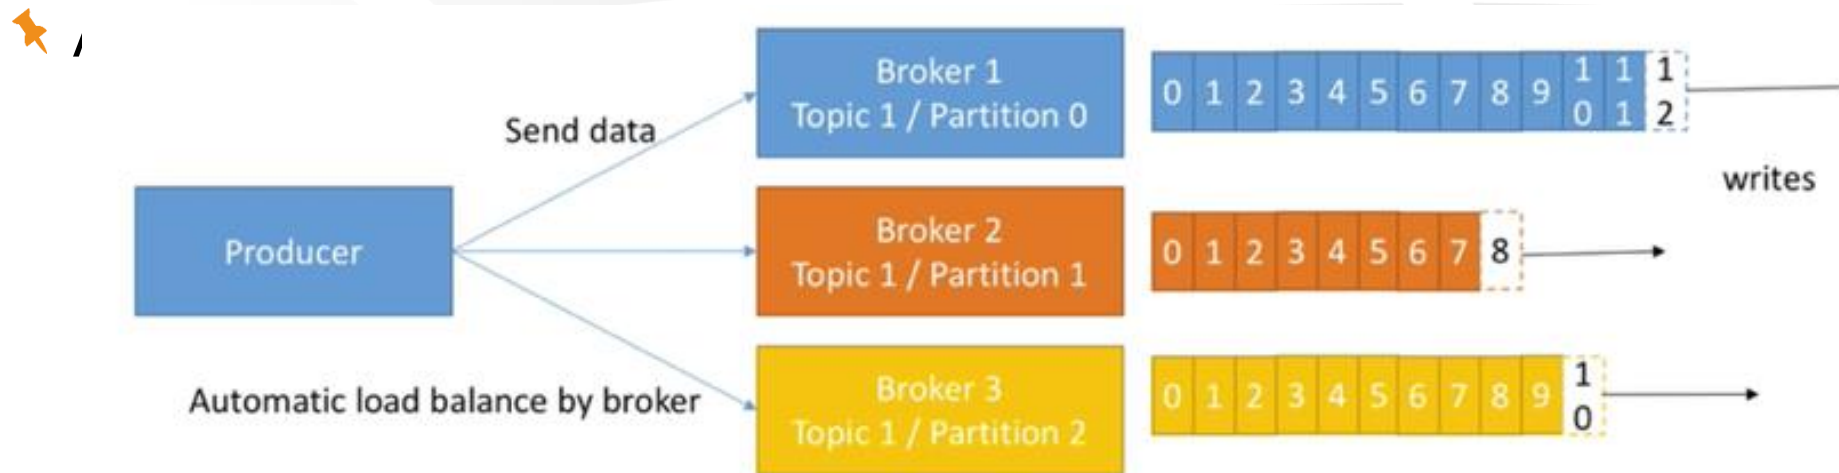
# Producers

- ✦ Producers write data to the partitions
- ✦ In order to send data a producer needs to know:
  - ✦ Topic name
  - ✦ At least one broker to connect to
  - ✦ Kafka will do the routing to the right brokers




# Producers

- ✦ Producers can choose to receive acknowledgment of data writes:
  - ✦ Acks = 0 : No waiting for acknowledgment ( possible data loss)
  - ✦ Acks = 1 : Producer will wait for leader acknowledgment (limited data loss)



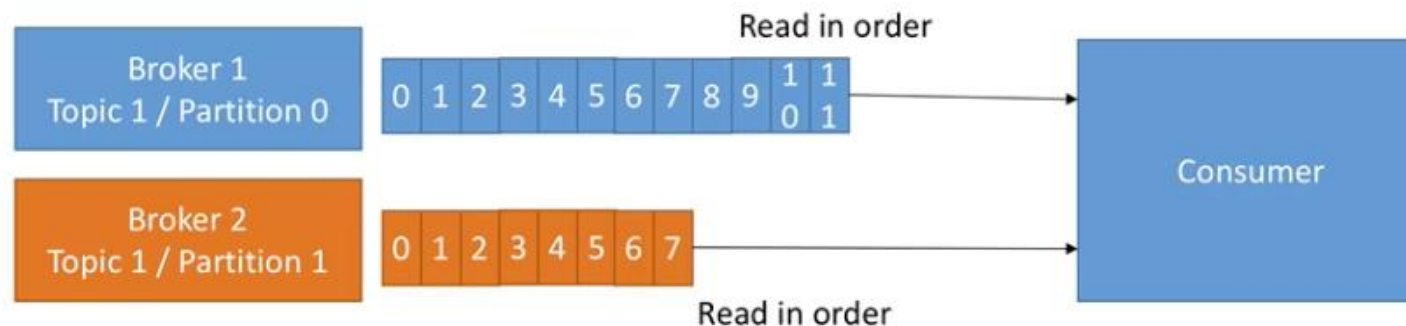
# Producers Message Keys

- ✦ Producers can choose to send a Key with a message
- ✦ When a key is sent, the producer has the guarantee that all message with that key will always go to the same partition.
- ✦ This enables Ordering for a specific Key !



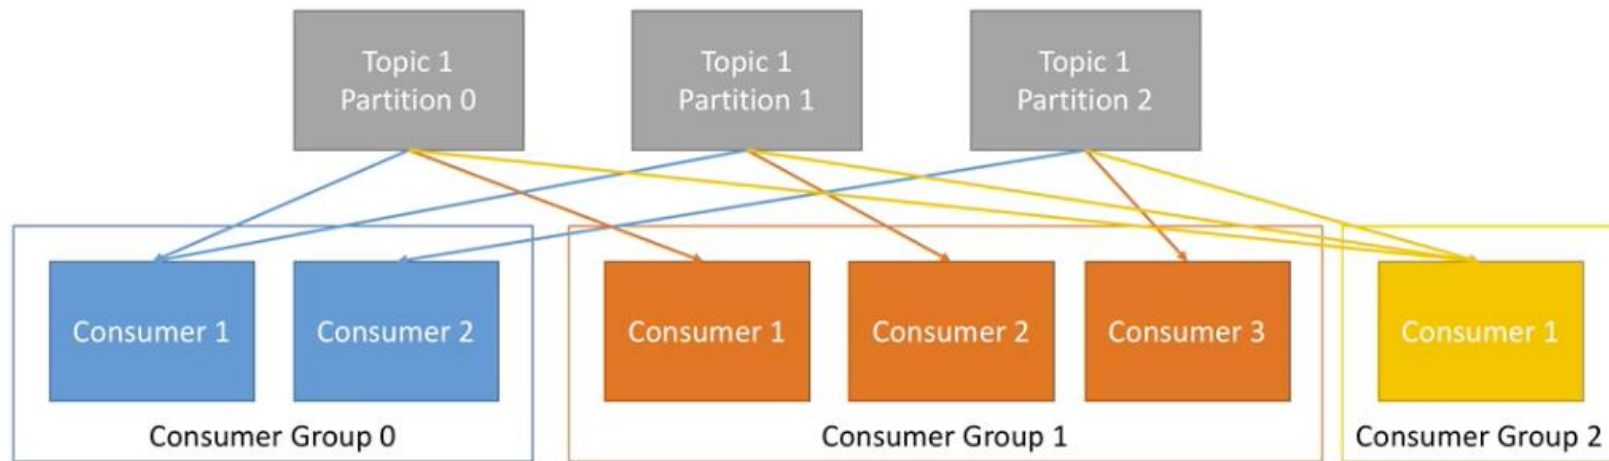
# Consumers

- ✦ Consumers read data from a partition
- ✦ Consumer need to specify :
  - ✦ The topic name
  - ✦ One broker from the broker list to connect to -> Kafka will take care of pulling the right data from the right broker.
- ✦ Data is read in-order for each partition.



# Consumer Groups

- ✦ Consumer read data within a specific consumer group
- ✦ Each consumer read from exclusive partitions
- ✦ You cannot have more consumers than partitions (otherwise some will become inactive)



# Consumer Offsets

- ✦ Kafka Stores the offsets at which a consumer group has been reading.
- ✦ The offset commits are store in a dedicated topic called `"__consumer_offsets"`
  - ✦ When consumer has processed the data successfully It should commit the offset
  - ✦ If a consumer process "crashes", it will be able to read back from where it left.



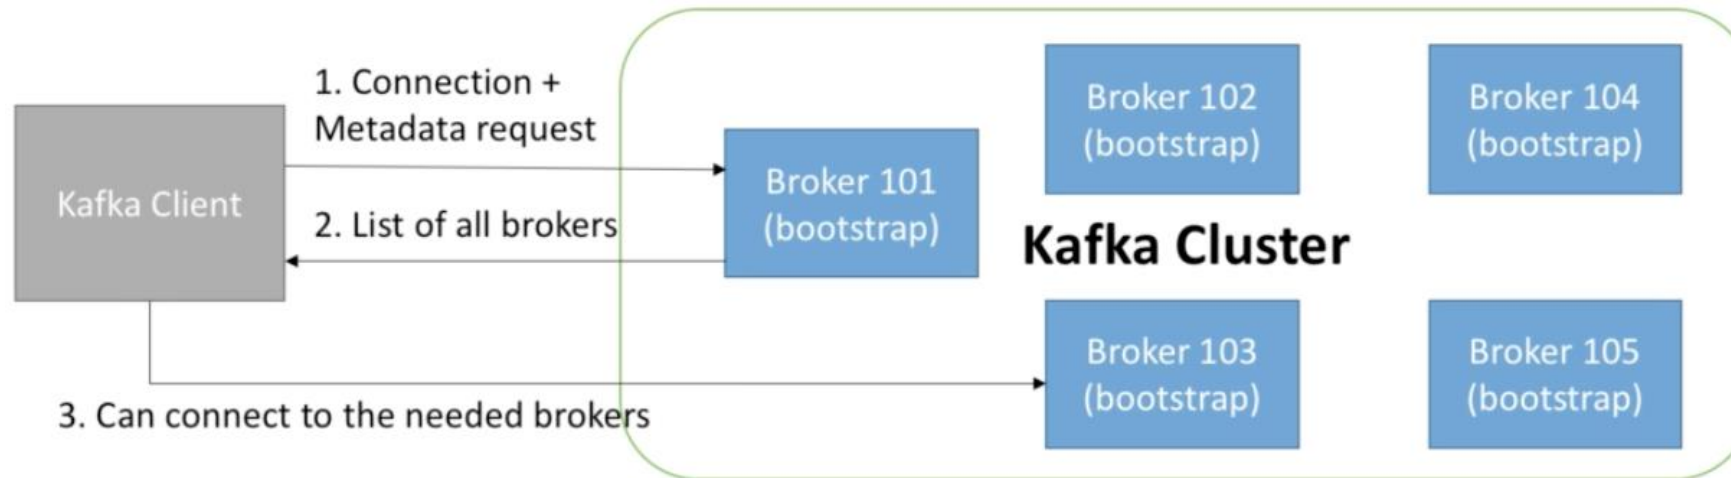
# Delivery Semantics for consumers

- ✦ Consumers can choose when to commit offsets.
  - ✦ There are 3 delivery semantics options:
    - ✦ **At Most Once**
      - Offsets are committed as soon as the message is received.
      - If the processing goes wrong, the message will not be read again.
    - ✦ **At least Once (preferred)**
      - Offsets are committed after message is processed.
      - If the processing goes wrong, the message will be read again.
      - Make sure processing the same message twice won't impact your system.
    - ✦ **Exactly Once**
      - Can only be achieved from Kafka to Kafka using K-Streams API.
-



# Kafka Broker Discovery

- ✦ Every kafka broker is called a “bootstrap server”
- ✦ That mean you only need to connect to one broker and you will be connected to the entire cluster
- ✦ Each broker knows about all the brokers, topics and partitions





# Zookeeper

- ✦ Zookeeper manages brokers (keeps a list of them)
- ✦ Zookeeper helps in performing leader election for partitions
- ✦ Zookeeper sends notifications to kafka in case of changes (e.g. new topic, broker dies, broker recovery, delete topics, etc.)
- ✦ **Kafka Can't work without Zookeeper**
- ✦ By design works with odd number of servers (3,5,7)
- ✦ Zookeeper has a leader (handles writes), the rest are followers
- ✦ Zookeeper does not store consumer offset (with Kafka > v0.10)

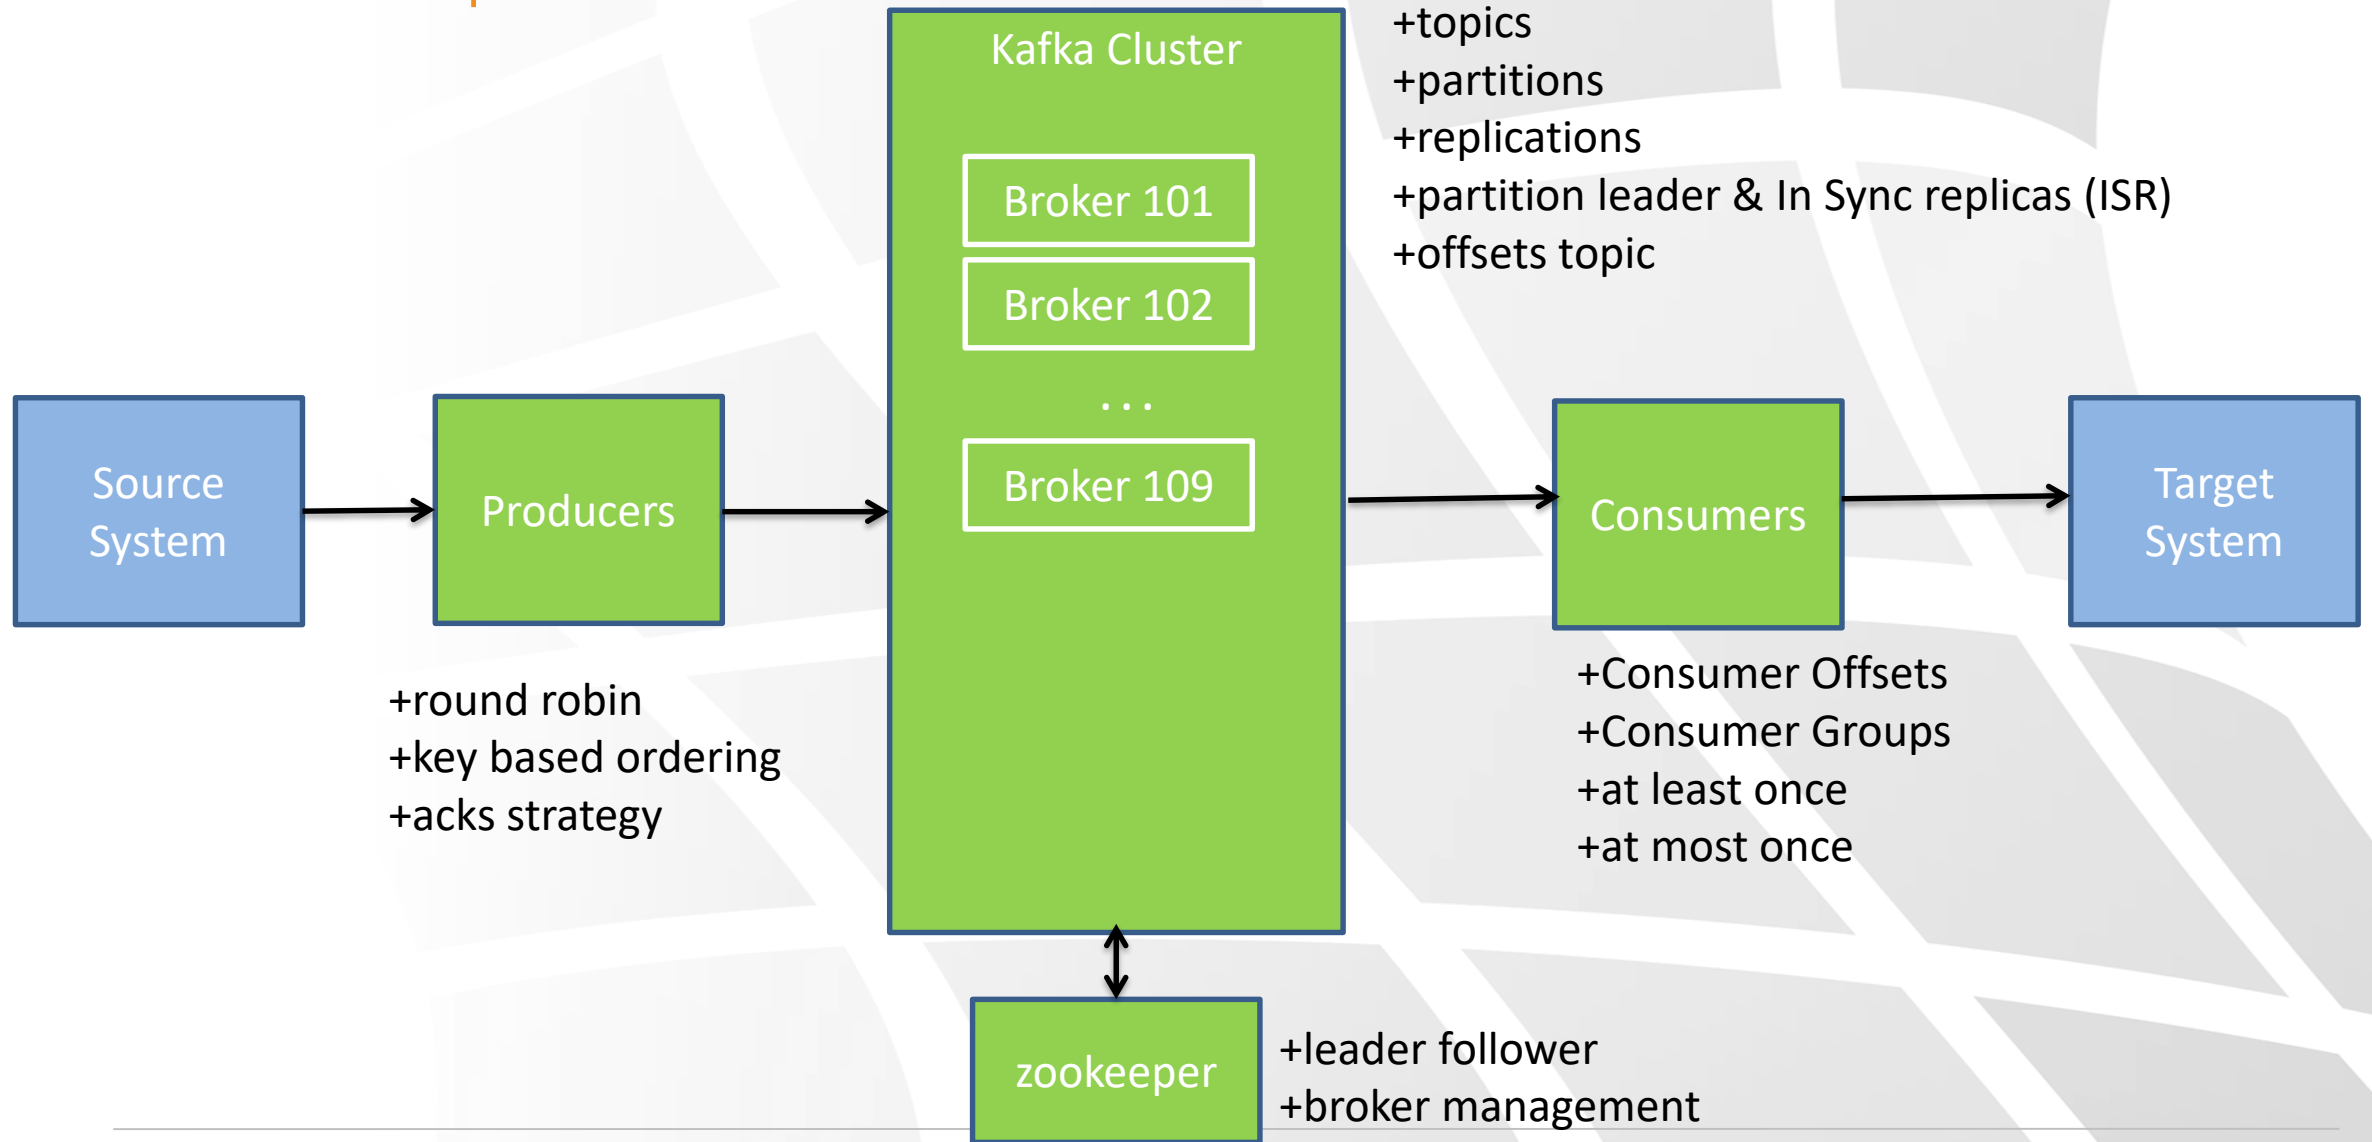


ZooKeeper

# Kafka Guarantees

- ✦ Messages are appended to a topic-partition in the order they are sent
  - ✦ Consumers read messages in the order stored in a topic-partition
  - ✦ With replication factor of  $N$ , producers and consumers can tolerate up to  $N-1$  brokers being down.
    - ✦ Replication factor of 3 :
      - ✦ One broker can be down for maintenance
      - ✦ Another broker can be down unexpectedly
  - ✦ As long as the number of partitions does not change, the same key will always go to the same topic.
-

# Roundup



# Summary

- ✦ Kafka Data is organized into streams of data called "Topics"
  - ✦ Each topic is organized within partitions.
  - ✦ A producer sends message with a key that guarantees an assignment to a specific partition.
  - ✦ Consumers are organized into consumer groups
  - ✦ Each consumer has a exclusive sets of partition he is assigned to.
  - ✦ Group offset are maintained by committing an offset by the consumer.
-



## Module 02: Installing Kafka



# Install Process

- ✦ Typically Kafka Installation will contain the following components :
  - ✦ Kafka Brokers
  - ✦ Zookeeper
  - ✦ Schema Registry (Optional )
- ✦ You can install each of those separately but we will install Docker in order to enable fast deployment of all the components in one container

# Manual Installation

- ✦ Download the Kafka binaries from <https://kafka.apache.org/downloads>
- ✦ The bin folder contains shell scripts for windows and Linux
- ✦ You need to run first zookeeper (from the root folder):

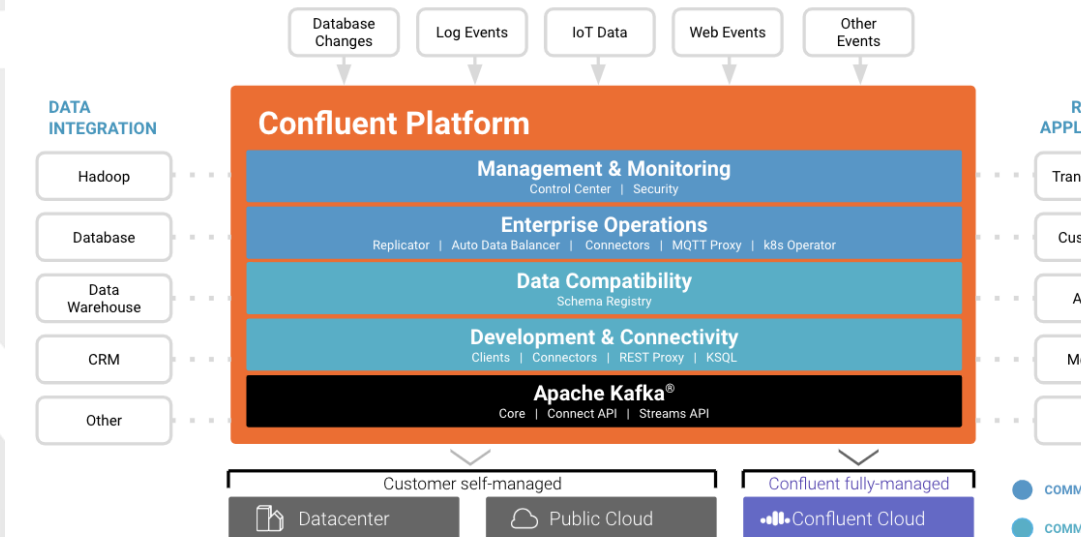
```
$ bin/zookeeper-server.start.sh config/zookeeper.properties
```

- ✦ Than start Kafka brokers:

```
$ bin/kafka-server.start.sh config/server.properties
```

# Confluent Platform

- ✦ Was built from the creators of Apache-kafka to enable enterprise solution for kafka
- ✦ Confluent Platform contains all of Kafka Echo system
  - ✦ Brokers
  - ✦ Kafka Streams + KSQL
  - ✦ Data Connectors
  - ✦ Client libraries (C++/python/ Go / Java/ .
  - ✦ Schema registry and the avro format
  - ✦ Monitoring, Control Center





# Dev Environment (fast-data-dev image for Docker)

# Docker for Mac >= 1.12, Linux, Docker for Windows 10

```
$docker run --rm -it \  
-p 2181:2181 -p 3030:3030 -p 8081:8081 \  
-p 8082:8082 -p 8083:8083 -p 9092:9092 \  
-e ADV_HOST=127.0.0.1 \  
landoop/fast-data-dev
```

# Kafka Endpoints

- ✦ A typical Kafka installation will use the following ports for several kafka services :
    - ✦ Kafka Brokers : port 9092 -> \*
    - ✦ Zookeeper : port 2181
    - ✦ Schema Registry : 8081
    - ✦ Rest Proxy : 8082
  - ✦ These can be configured in `zookeeper.properties/server.properties`
  - ✦ During this course we assume a kafka installation in at `~/kafka`
  - ✦ Our executables will be at `~/kafka/bin`
-

Starting Landoop fast-data-dev environment

# Demo



# Lab 01: Interacting with Ladoop Kafka Environment

## Lab



<https://github.com/selagroup/KafkaWorkshop/blob/master/Lab-01.md>



## Module 03: Kafka CLI

Kafka Workshop



# Agenda

- ✦ Kafka CLI
- ✦ List All Topics
- ✦ Create and remove a topic
- ✦ Kafka Producer
- ✦ Kafka Consumer
- ✦ Consumer Groups

# The Kafka CLI

- ✦ CLI tool can be accessed by downloading the Kafka Binaries:  
<https://kafka.apache.org/downloads>
- ✦ Make sure to add the /bin folder to your PATH
- ✦ If you Want to work with a predefined docker image :
  - ✦ `docker run --rm -it --net=host landoop/fast-data-dev bash`

# \$ kafka-topics

```
$ kafka-topics.sh --zookeeper 127.0.0.1:2181 --list
```

```
$ kafka-topics.sh --zookeeper 127.0.0.1:2181 --topic myTopic --partitions 3 --replication.factor 2
```

```
$ kafka-topics.sh --zookeeper 127.0.0.1:2181 --topic myTopic --describe
```

- ✦ Create/Delete/Describe or Change a Topic
- ✦ Need to specify the zookeeper address
- ✦ Flags:
  - ✦ --list : list all topics
  - ✦ --create --topic topic\_name
  - ✦ --partitions
  - ✦ --replication-factor
  - ✦ --describe



# \$ kafka-console-producer

```
$ kafka-console-producer --broker-list 127.0.0.1:9092 --topic myTopic
```

- ✦ Publish data to kafka topic using a strings as values and keys
- ✦ When no key is specified a random key is generated.
- ✦ If The topic does exists, Kafka will produce an error ,but will create the topic.

# \$ kafka-console-consumer

```
$ kafka-console-consumer --bootstrap-server 127.0.0.1:9092 --topic myTopic
```

```
$ kafka-console-consumer --bootstrap-server 127.0.0.1:9092 --topic myTopic --from-beginning
```

- ✦ Reads data from a default Consumer group into the standard output
- ✦ Note the --bootstrap-server flag vrs --broker-list flag
- ✦ Default behavior will read from the latest committed offset
- ✦ Read options includes
  - ✦ --from-beginning
  - ✦ --group group.id=group1
  - ✦ --max-messages 100

# Consumer Groups

⚡ What happens when you run the consumer command over and over again on the same consumer group ?

```
$ kafka-console-consumer --bootstrap-server 127.0.0.1:9092 --topic myTopic --group grp1 --from-beginning
```

# \$ kafka-consumer-groups

```
$ kafka-consumer-groups --bootstrap-server 127.0.0.1:9092 --describe --group grp1
```

```
$ kafka-consumer-groups --bootstrap-server 127.0.0.1:9092 --list
```

✦ Lists all consumer groups, describe a consumer group ,delete and reset.

✦ --reset-offsets allow resetting the consumer group offset

✦ --to-date-time

✦ --by-period

✦ --to-earliest

✦ --to-latest

✦ --shift-by

```
$ kafka-consumer-groups --bootstrap-server 127.0.0.1:9092 --reset-offsets --shift-by --2 --group grp1 --execute
```

# Producer/Consumer with a custom 'Key'


```
$ kafka-console-producer --broker-list 127.0.0.1:9092 --topic myTopic \  
--property "parse.key=true" \  
--property "key.separator=:"
```

- ✦ It is possible to specify a dedicated key when publishing a message
- ✦ Kafka will hash all keys into the same partition Id

```
$ kafka-console-consumer --bootstrap-server 127.0.0.1:9092 --topic myTopic \  
--property "print.key=true" \  
--property "key.separator=:"
```

# Kafka Connect with the CLI Tool (Module 7)

## Kafka Connect



## Lab 02: Basic commands

# Lab



<https://github.com/selagroup/KafkaWorkshop/blob/master/Lab-02.md>

# Questions







# Module 04: Programmatic API

Kafka Workshop



# Agenda

- ✦ The Kafka SDK For Java
- ✦ Producer API
- ✦ Consumer API
- ✦ Consumer Groups and rebalance
- ✦ Using Configuration files
- ✦ Using Custom Serializers

# Producer API

✦ The Producer API allows applications to send streams of data to Kafka  
<https://kafka.apache.org/documentation/#producerapi>

✦ Maven Config :

```
<dependency>  
  <groupId>org.apache.kafka</groupId>  
  <artifactId>kafka-clients</artifactId>  
  <version>2.1.0</version>  
</dependency>
```

# Producer Config

- `bootstrap.servers` : a list of one or more of brokers
- `key.serializer/value.serializer` : the serialization classes for the keys and values
- `acks` (0/1/all) : request for either the leader or its replicants to send ack after sending.
- `retries` : In case of failure this field controls the retry count. Effects Ordering !
- `linger.ms` : batches all the request that arrive within the time frame and send them as single batch. This may effect latency, but reduces number of

# Building and using a Producer

```
public static KafkaProducer<String,String> BuildProducer(String brokers)
{
    Properties properties = new Properties();
    properties.setProperty("bootstrap.servers",brokers);
    properties.setProperty("key.serializer", StringSerializer.class.getName());
    properties.setProperty("value.serializer", StringSerializer.class.getName());

    properties.setProperty("acks", "1");
    properties.setProperty("retries","3");
    properties.setProperty("linger.ms","1");

    return new KafkaProducer<>(properties);
}
```

```
producer.send(new ProducerRecord<String, String>(topic , sentence));
```

# Consumer Config

- `bootstrap.servers` : a list of one or more of brokers
- `key.deserializer/value.deserializer` : the deserialization classes for the keys and values
- `enable.auto.commit` : consumer will commit the offset automatically in the background.
- `group.id` : the consumer group id.
- `auto.offset.reset` : behavior when there is no offset (earliest/latest/none)

# Building and using a Consumer

```
public static KafkaConsumer<String,String> BuildConsumer(String brokers, String groupId)
{
    Properties properties = new Properties();
    properties.setProperty("bootstrap.servers",brokers);
    properties.setProperty("bootstrap.servers",brokers);
    properties.setProperty("key.deserializer", StringDeserializer.class.getName());
    properties.setProperty("value.deserializer", StringDeserializer.class.getName());

    properties.setProperty("enable.auto.commit","true");
    properties.setProperty("group.id",groupId);
    properties.setProperty("auto.offset.reset","earliest");

    KafkaConsumer<String,String> consumer = new KafkaConsumer<~>(properties);
    return consumer;
}
```

```
consumer.subscribe(Arrays.asList(topic));
ConsumerRecords<String,String> records = consumer.poll(Duration.ofMillis(500));
```

# Simple Producer and Consumer

## Demo





# Working with config files

- It is advised to work with configuration files to describe the various kafka configuration.
- A recommended pattern is to use the HCON file format with the *TypeSafe* maven repository to describe the configuration

```
<dependency>  
  <groupId>com.typesafe</groupId>  
  <artifactId>config</artifactId>  
  <version>1.3.3</version>  
</dependency>
```

```
{  
  kafka {  
    topic-name: "my-topic",  
    producer {  
      bootstrap.servers = "localhost:9092",  
      key.serializer = org.apache.kafka.common.serialization.StringSerializer,  
      value.serializer = org.apache.kafka.common.serialization.StringSerializer,  
      acks = 1,  
      retries = 3,  
      linger.ms = 1  
    },  
    consumer {  
      bootstrap.servers = "localhost:9092",  
      key.deserializer = org.apache.kafka.common.serialization.StringDeserializer,  
      value.deserializer = org.apache.kafka.common.serialization.StringDeserializer,  
      enable.auto.commit = true,  
      group.id = grp1,  
      auto.offset.reset = "earliest"  
    }  
  }  
}
```

# Working with consumer groups and re-balance

- When we ran the several consumer processes on the same consumer group a rebalance phase will occur.
- Look for the “AbstractCoordinator” logs that describe the re-assignments of partitions across consumers , e.g.
  - Revoking previously assigned partitions[ my-topic-0,my-topic-1]
  - Setting newly assigned partitions[ my-topic-0]

```
s.AbstractCoordinator - [Consumer clientId=consumer-1, groupId=grp12] (Re-)joining group
s.AbstractCoordinator - [Consumer clientId=consumer-1, groupId=grp12] Successfully joined group with generation 36
s.ConsumerCoordinator - [Consumer clientId=consumer-1, groupId=grp12] Setting newly assigned partitions [my-topic-0, my-topic-1]
s.AbstractCoordinator - [Consumer clientId=consumer-1, groupId=grp12] Attempt to heartbeat failed since group is rebalancing
s.ConsumerCoordinator - [Consumer clientId=consumer-1, groupId=grp12] Revoking previously assigned partitions [my-topic-0, my-topic-1]
s.AbstractCoordinator - [Consumer clientId=consumer-1, groupId=grp12] (Re-)joining group
s.AbstractCoordinator - [Consumer clientId=consumer-1, groupId=grp12] Successfully joined group with generation 37
s.ConsumerCoordinator - [Consumer clientId=consumer-1, groupId=grp12] Setting newly assigned partitions [my-topic-0]
sumer - got record #1 partition:0 offset : 1683
```

# Interacting with Zookeeper

- ✦ In case you want to create/delete and manage topics, you will need to interact with zookeeper itself .
- ✦ The AdminZkClient is a utility class that can assist in those tasks :
  - ✦ Create, delete and modify a topic
  - ✦ Add partitions to a topic
  - ✦ List all topics and their configuration
  - ✦ Change the broker configuration

```
private static void listTopics() {  
    KafkaZkClient zkClient = KafkaZkClient.apply(zookeeperHost, isSucre, sessionTimeoutMs,  
        connectionTimeoutMs, maxInFlightRequests, time, metricGroup, metricType);  
  
    AdminZkClient adminZkClient = new AdminZkClient(zkClient);  
    Map<String, Properties> topics = adminZkClient.getAllTopicConfigs();  
    scala.collection.Iterator iter = topics.keysIterator();  
    while (iter.hasNext()) {  
        System.out.println(iter.next());  
    }  
}
```

# Assign and Seek

- ✦ Sometimes it is desirable to read from a specific topic or even a certain offset
    - ✦ Playback “bad” input stream that was streamed incorrectly
    - ✦ Logging and tracing
    - ✦ Performance testing
  - ✦ The class “*TopicPartition*” determines the partition from which to read from a specific topic
  - ✦ The Assign and Seek method assist in seeking
  - ✦ No consumer group is involved in this process.
-

Assign and Seek

# Demo



# Lab 03: Java API and Re-balance

## Lab



<https://github.com/selagroup/KafkaWorkshop/blob/master/Lab-03.md>

# Questions





## Module 05: Advanced Programming & Delivery Semantics





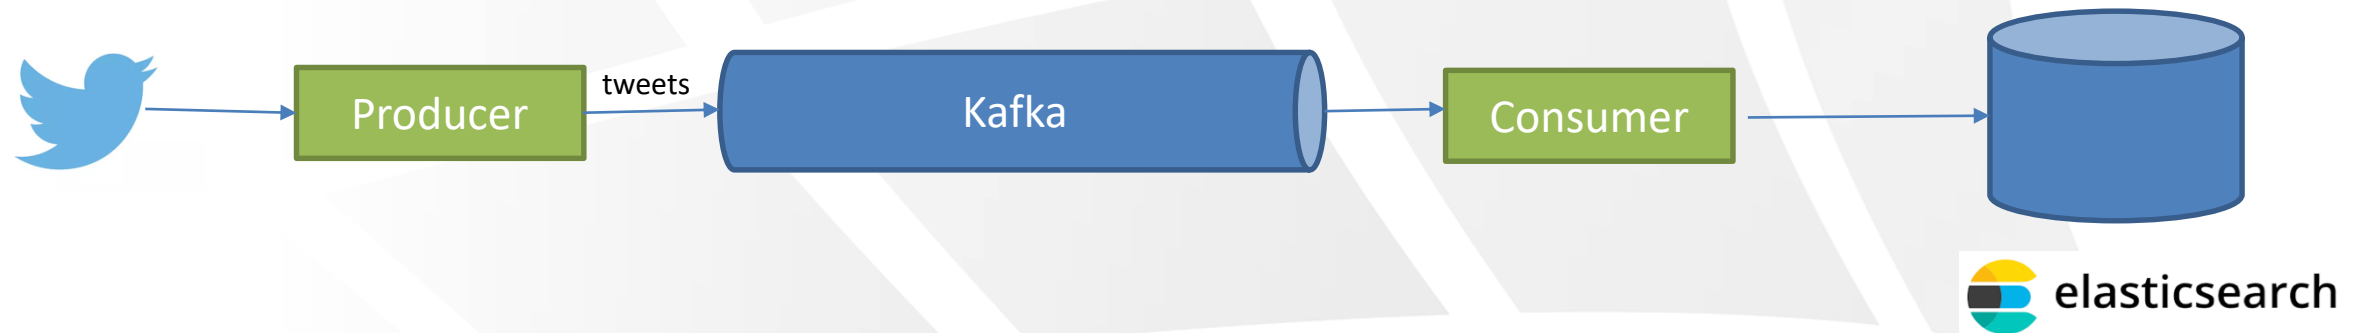
# Agenda

- ✦ Advance sample app with Twitter API integration
- ✦ High Throughput and Safe Producer
- ✦ Message Compression
- ✦ Kafka Elastic Search Consumer + Advance Config
- ✦ Idempotent Producer and Consumer

---

- ✦ **Delivery Sematcis**

# Twitter to Elastic Search Demo



- ✦ We will build a highly scalable twitter collector using kafka API Concepts:
- ✦ Producer will be reliable (retries mechanism, compression, batch)
- ✦ Consumer delivery semantics ("at-most-once" / "at-least-once" / etc)
- ✦ Offset groups reset in case we want to "replay"

# Producers Acks = 0 (no acks)

- ✦ No response is requested
- ✦ If the broker goes offline or an exception is raised, we don't know and we loose data



- ✦ Useful for cases where data loss is acceptable :
  - ✦ Metrics collection
  - ✦ Log collection

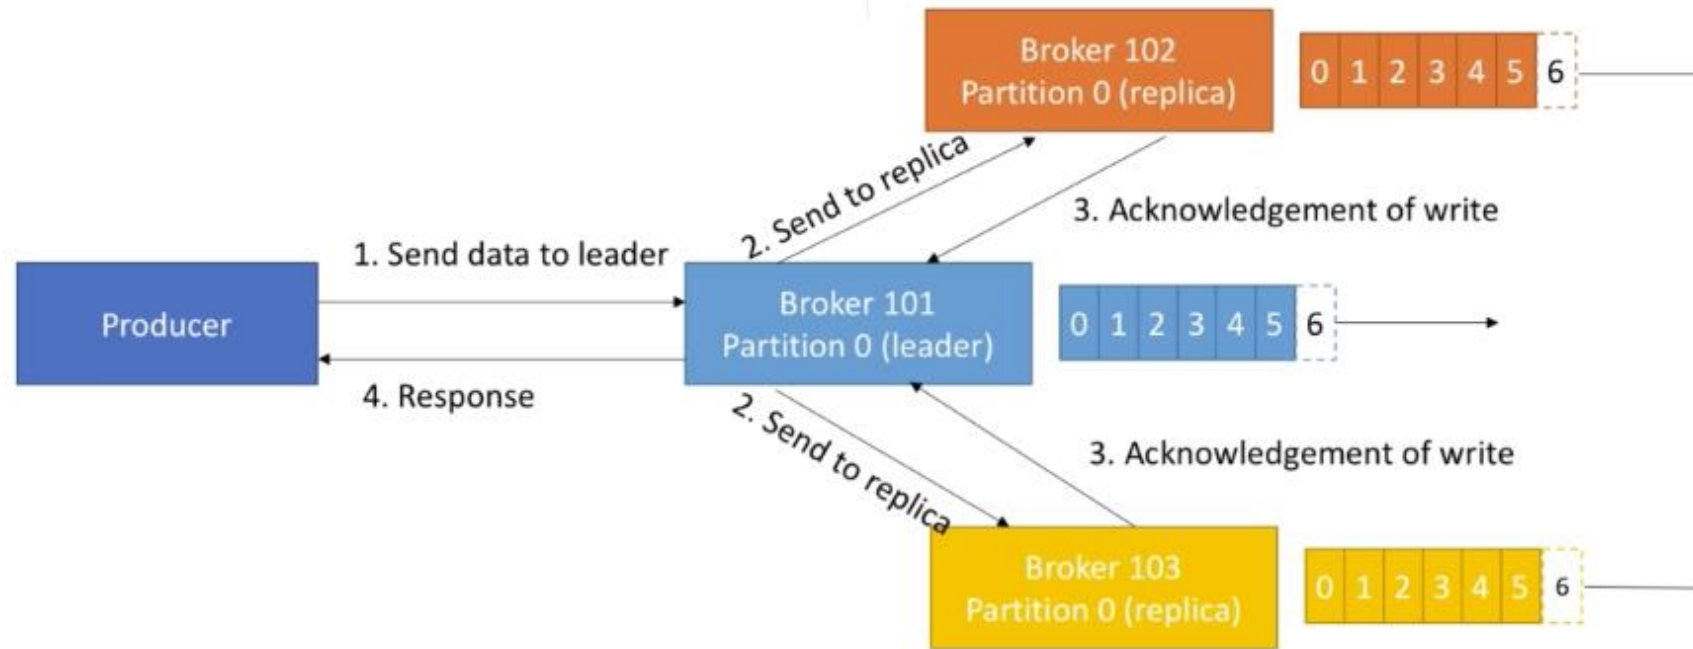
# Producers Acks = 1 (leader acks)

- ⚡ Leader response is requested but replication is not guaranteed (background)
- ⚡ If the leader brokers goes offline before replication was completed we have data loss.



# Producers Acks = 2 (replicas acks)

✦ Leader + Replicas ack Requested



✦ Added latency and safety

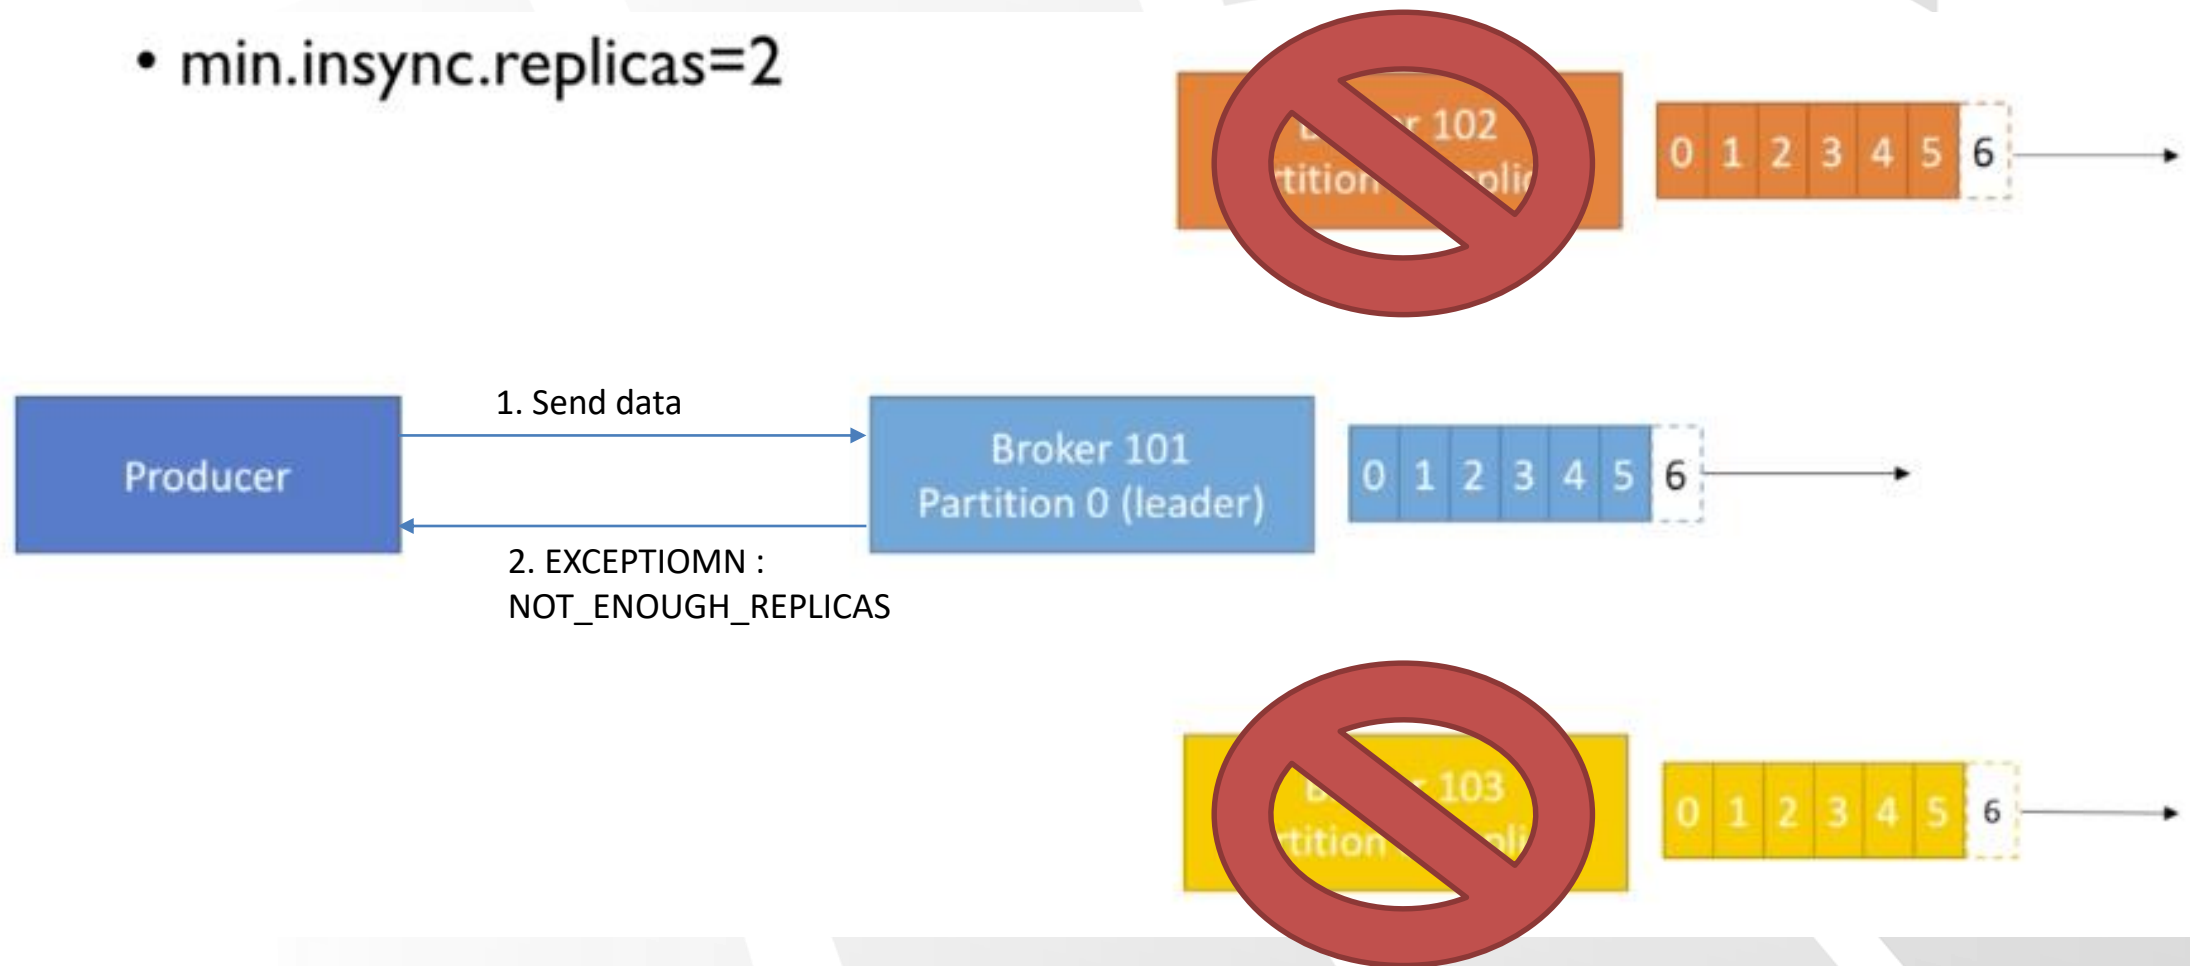
✦ Necessary if you don't want to loose data

# Producers Acks =2 (replicas acks)

- ✦ Acks=all must be used in conjunction with `min.insync.replicas=2`
- ✦ `min.insync.replicas=2` can be set at the broker or topic level (override)
- ✦ `min.insync.replicas=2` implies that at least 2 brokers that are ISR (including the leader ) must respond that they have the data.
- ✦ That means that if you use `replication.factor=3,min.insync.replicas=2 , acks=all`, you can only tolerate 1 broker going down, otherwise the producer will receive an Exception on send

# Producers Acks = 2 (replicas acks)

- `min.insync.replicas=2`



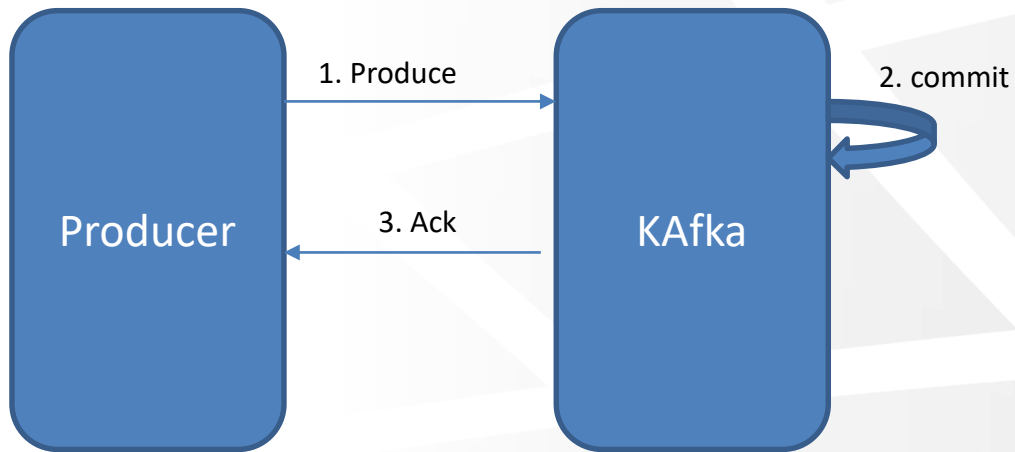
# Producer retries

- ✦ In case of transient failures, developers are expected to handle exceptions, otherwise the data will be lost.
- ✦ Example of transient exceptions :
  - ✦ `NotEnoughReplicasException`
- ✦ There is a “retries” setting:
  - ✦ Default is 0
  - ✦ Can increase to a high number e.g. `Integer.MAX_VALUE`
  - ✦ In case of retries there is a chance a message will be sent out of order
  - ✦ We can control parallelism with `max.in.flight.requests.per.connection` (default is 5)
    - ✦ Set it to 1 ensures ordering.
    - ✦ There is a better solution in Kafka > 1.0.0 !

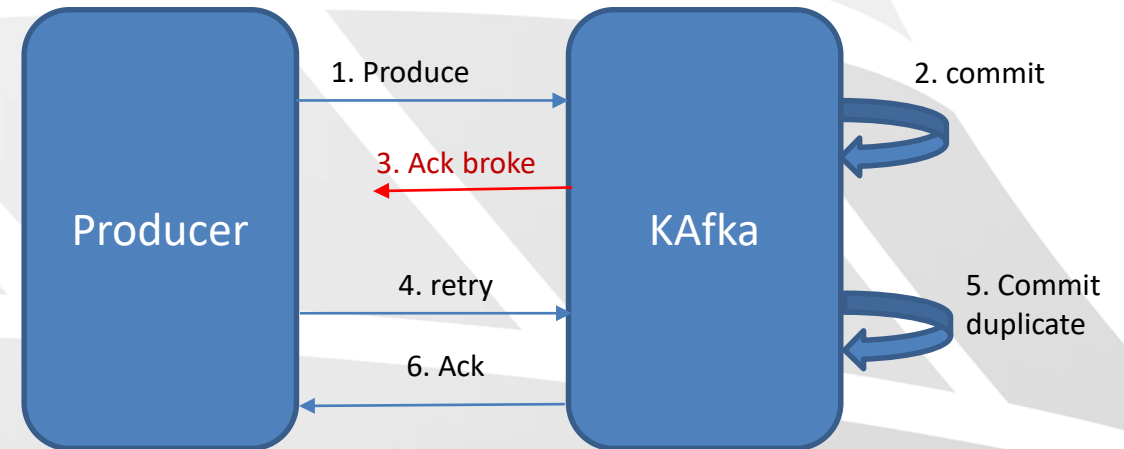


# Idempotent Producer – The problem

“Good” Request

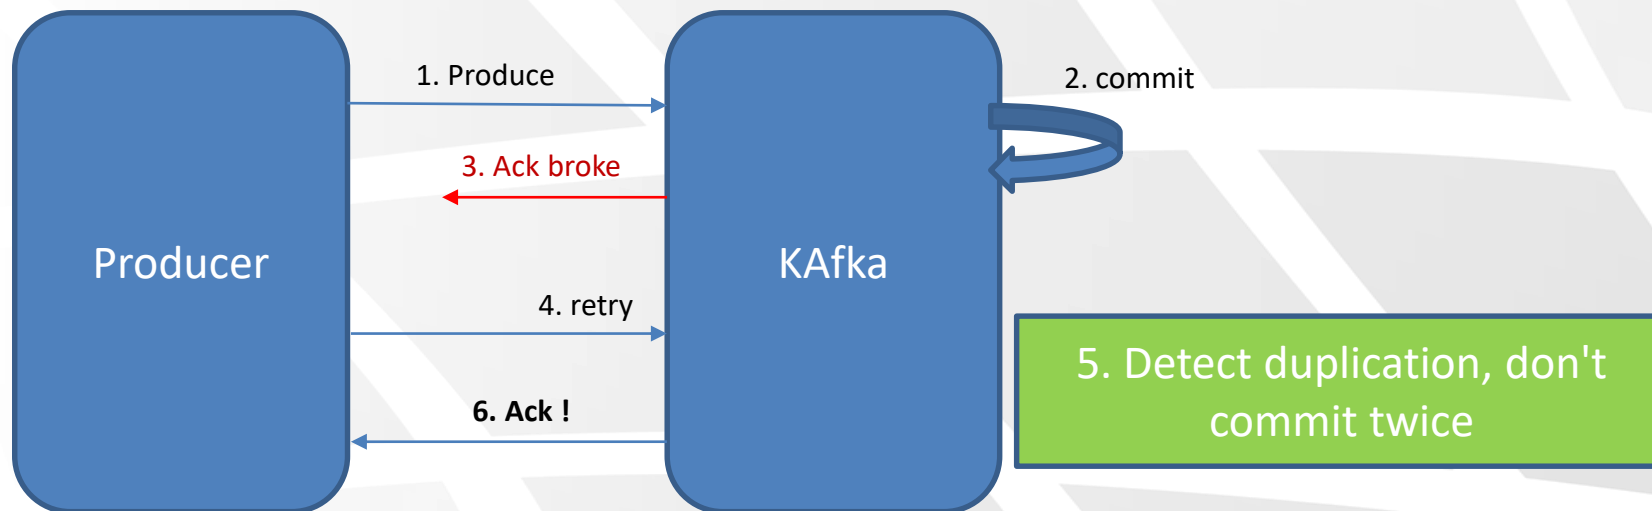


“Duplicate message” Request



# Idempotent Producer – Solution (Kafka >=0.11)

✦ In Kafka >=0.11 , you can define an “idempotent producer” which won’t introduce duplicate on network failure



# Idempotent Producer

- ✦ Idempotent producer are great to guarantee a safe and stable pipeline
- ✦ They must come with :
  - ✦ `retries = Integer.MAX_VALUE`
  - ✦ `max.in.flight.requests = 1` (Kafka < 1.1)
  - ✦ `max.in.flight.requests = 5` (Kafka > 1.1 – higher performance)
  - ✦ `acks = all`
- ✦ We need to set
  - ✦ **`producerProperties.put("enable.idempotence", true)`**

# Safe Producer – Summary & Demo

## ✦ For Kafka < 0.11

- ✦ `acks = all` (producer level)
- ✦ `min.insync.replicas=2` (topic/broker level)
- ✦ `retries = MAX_INT` (producer level)
- ✦ `max.inflight.requests.per.connection=1` (producer level)

## ✦ For Kafka >= 0.11

- ✦ `enable.idempotence=true`
- ✦ `min.insync.replicas=2`
- ✦ Implies:
  - ✦ `Acks = all` , `retries= MAX_INT` , `max.inflight.requests.per.connection=5`
  - ✦ Order is maintained and performance is high !



Running a safe producer might impact throughput, always test for your case !

# Using Safe / Idempotent Producer

## Demo



# Message Compression

- ✦ Producer usually sends data that is text based, for example Json
- ✦ In such case It is important to enable compression to the producer.
- ✦ This does not effect the consumer and/or the broker
- ✦ `compression.type` can be:
  - ✦ `none`
  - ✦ `gzip`
  - ✦ `lz4`
  - ✦ `snappy`
- ✦ The bigger the batch is , the better the compression is .
- ✦ Benchmark : <https://blog.cloudflare.com/squeezing-the-firehose/>

# Compression benefits

## ✦ Advantage of using compression:

- ✦ Much smaller producer request size ( compression ration up to x4 !)
- ✦ Faster to transfer the data over the network
- ✦ Better throughput
- ✦ Better Disk utilization

## ✦ Disadvantages:

- ✦ Producers commit some CPU cycles is used for the compression
- ✦ Consumers commit some CPU cycles to decompress data

## ✦ Overall :

- ✦ Consider testing for compression against snappy and lz4 for speed/ratio
- ✦ Always use compression in production
- ✦ Consider tweaking [linger.ms](#) and [batch.size](#) to have bigger batches

# linger.ms and batch size

- ✦ By default Kafka tries to send messages as soon as possible
  - ✦ It will have up to 5 queues in-flight, meaning up to 5 messages are sent concurrently.
  - ✦ In the background other messages are already being batched while they wait.
  - ✦ This allows kafka to increase throughput with low latency !
  - ✦ Batches have higher compression rate.
- ✦ **linger.ms**
  - ✦ The number of milliseconds a producer is willing to wait before sending a batch
    - ✦ default : 0
    - ✦ By providing some lag (e.g `linger.ms = 5`) ,we increase the chance of batching
    - ✦ If batch is full (see `batch.size`) before the end of **linger.ms** ,it is being sent immediately.
- ✦ **batch.size :**
  - ✦ Maximum number of bytes that will be included in a batch. default is 16KB
  - ✦ Increasing to 32K or 64K can improve compression, hence throughput
  - ✦ Any message that is bigger than the batch size will not be batched
  - ✦ A batch is allocated per partition ! **Setting it to a large number may drain your memory !**



# High Throughput Producer

## Demo




- Adding snappy
- Increase batch size to 32K
- Add delay of linger.ms = 20ms

# How Keys are hashed

- ✦ By default, your keys are hashed using the “murmur2” algorithm
- ✦ It is preferred not to override the behavior of the partitioner, but it is possible to do so ([partitioner class](#))
- ✦ The formula is  
`targetPartitioner = Utils.abs(Utils.murmur2(record.key()))%numPartitions;`
- ✦ Same key will always go to the same partition, but adding partition will break keying.

# Consumer And Elastic Search



# Create a new Elastic Search Cluster (For Dev)

- ✦ The leader search engine platform for distributed ,multitenant full text-search.
- ✦ Json-based docs with backed schema
- ✦ Create a free Elastic Search cluster at :  
<https://bonsai.io/>
- ✦ Start with exploring the cluster :  
<https://www.elastic.co/guide/en/elasticsearch/reference/current/getting-started-explore.html>
- ✦ Create and list Indexes
  - ✦ PUT /twitter
  - ✦ GET /\_cat/indices



Using an ElasticSearch Cluster and adding docs

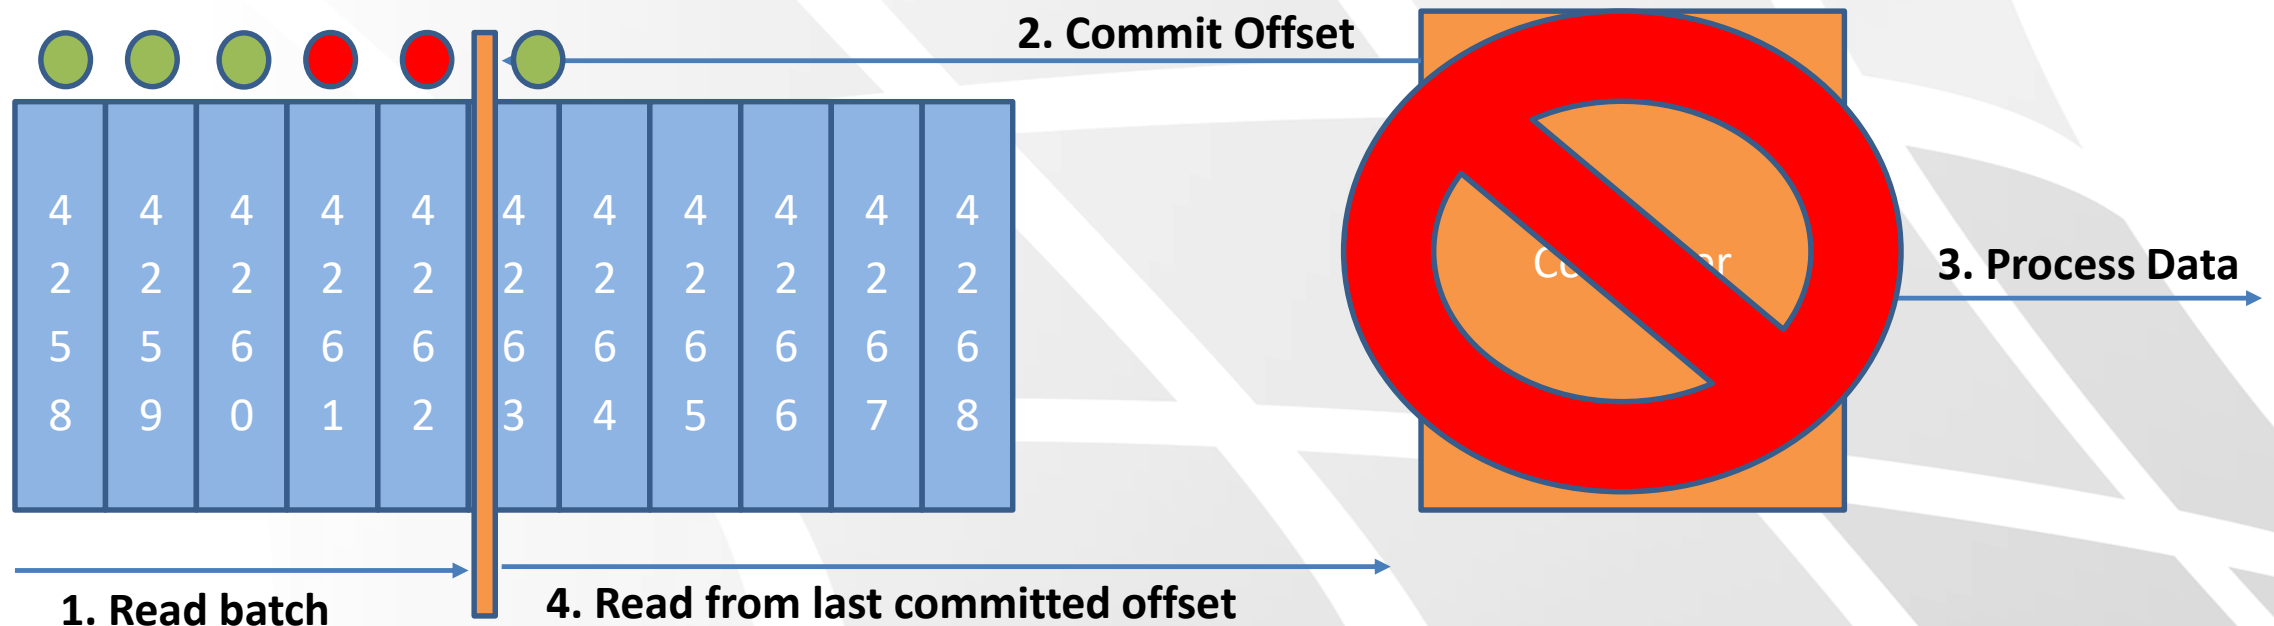
# Demo



# Delivery Semantics for Consumers – At most Once

## ✦ At Most Once

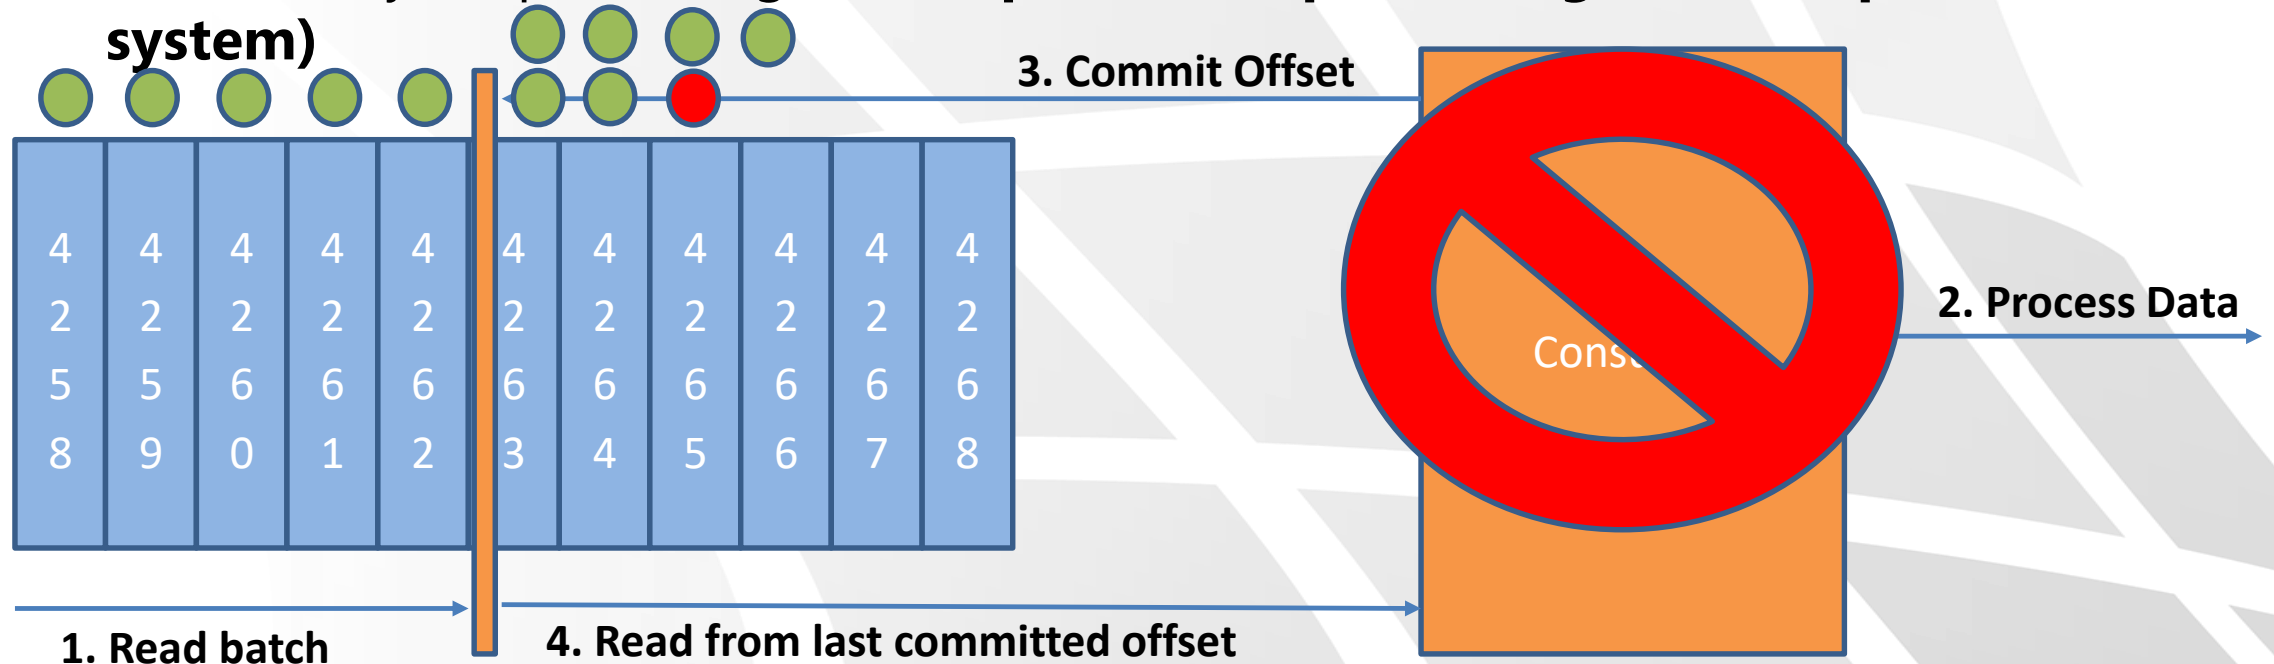
- ✦ Offsets are committed as soon as the message batch is received
- ✦ If the processing goes wrong (crashes) the message is lost .



# Delivery Semantics- At Least Once

## ✦ At Least Once

- ✦ Offsets are committed after the message is processed
- ✦ If the processing goes wrong (crashes) the message will be read again.
- ✦ Make sure your processing is **Idempotent (re-processing won't impact the system)**



# Delivery Semantics - Summary

## ✦ At most once

- ✦ Offsets are committed as soon as the message is received . Possible data loss.

## ✦ At least Once (default)

- ✦ Offset are committed after the message is processed. Something goes wrong same message can be read twice and processed twice. Make sure your processing is **Idempotent** (e.g.: Processing again the message won't impact your system).

## ✦ Exactly Once

- ✦ Only works between Kafka to Kafka Workflows using Kafka Streams API

**Conclusion** : for most application use **at-least-once** and make sure your processing is idempotent



# Make the Consumer Idempotent

## Demo



- Generating unique Id's to make producer idempotent

# Consumer Poll Behavior

- ✦ Kafka Consumers have a “poll” model, while other messaging bus systems has a “push” model
- ✦ This allows Kafka consumers to control where in the log they want to consume, how fast, and gives them the ability to reply.
- ✦ **Fetch.min.bytes** (default 1):
  - ✦ How much data you want to poll at least on each request
  - ✦ Helps increase throughput in the cost of latency
- ✦ **Max.poll.records** (default 500)
  - ✦ Controls how many records received per poll request
  - ✦ Increase If message are very small and you have lots of RAM
  - ✦ Good to monitor how many records are pulled per request.

# Consumer Poll Behavior

## ✦ `Max.partition.fetch.bytes`( default 1 MB):

- ✦ Maximum data returned by the broker per partition
- ✦ If you read from 100 partitions you will need a lot of memory

## ✦ `Fetch.max.bytes`(default 50MB):

- ✦ Maximum data returned for each fetch request (covers multiple partitions)
- ✦ The consumer performs multiple fetches in parallel

# Consumer Offset Commit Strategies

- ✦ There are two most common patterns for committing offsets in a consumer application
- ✦ (easy) `enable.auto.commit = true` + synchronous processing of batches

```
while(true){  
    List<Records> batch = consumer.poll(Duration.ofMillis(100))  
    doSomethingSynchronous(batch)  
}
```

- ✦ With auto-commit, offsets will be committed automatically for you at regular interval (`auto.commit.interval.ms = 5000` by default) whenever you call **.poll**
- ✦ If you don't use synchronous processing you will be in "at-most-once"

behavior

# Consumer Offset Commit Strategies – Cont.

- ✦ `enable.auto.commit` = false + synchronous processing of batches

```
while(true){  
    batch += consumer.poll(Duration.ofMillis(100))  
    if isReady(batch) {  
        doSomethingSynchronous(batch)  
        consumer.commitSync();  
    }  
}
```

- ✦ You control when you commit offsets and what's the condition for committing them
- ✦ Example : accumulating records to a buffer and then flushing the buffer to a database + committing offset then

# Commit strategies + max poll size

## Demo



- Controlling the commit strategy by allowing manual commit
- Configuring max poll size
- Use Elastic Search Batches

# Consumer Offset Reset Behavior

- ✦ When our app has a bug and cannot read from kafka since the consumer are down , we get by default 7 days of retention
- ✦ The behavior for the consumer then should be :
  - ✦ `auto.offset.reset` = latest - will read from the end of the log
  - ✦ `auto.offset.reset` = earliest - will read from the start of the log
  - ✦ `auto.offset.reset` = none - will throw exception if no offset is found
- ✦ Consumer offsets can be lost :
  - ✦ If a consumer hasn't read new data in 1 day ( kafka < 2.0)
  - ✦ If a consumer hasn't read new data in 7 day ( kafka >= 2.0)
  - ✦ This can be controlled by the broker settings `offset.retention.minutes`

# Querying Consumer Groups offsets

✦ We can query what is the current offset of each consumer group:

```
$ kafka-consumer-groups.sh --bootstrap-server 127.0.0.1:9092 --group grp --describe
```

✦ We can reset a consumer group :

```
$ kafka-consumer-groups.sh --bootstrap-server 127.0.0.1:9092 --group grp --reset-offsets  
--execute --to-earliest --topic tweets
```



# Replaying data for Consumers

- ✦ To replay data for a consumer group :
  - ✦ Take all the consumer from a specific group down
  - ✦ Use 'Kafka-Consume-groups' command to set the offset
  - ✦ Restart Consumers

## **Summary :**

- ✦ Set proper `data.retention.period` and `offset.retention.period`
- ✦ Make sure the `auto.reset.behavior` is what you want (auto/manual)
- ✦ Use replay capability in case of unexpected behavior

# Questions






## Module 06: Kafka Streams (Intro)



# What is Kafka Streams

⚡ **A data processing and transformation library** within Kafka



- ⚡ Standard Java Application
- ⚡ No need to create separate cluster
- ⚡ Highly scalable, elastic and fault-tolerant
- ⚡ Exactly Once capabilities

# What is stream processing ?

- ✦ Stream processing is really all about **Transformation on a** continuous stream of data
- ✦ Transformation are in forms of filters, maps, joins and aggregations
- ✦ We can divide stream processing into two categories:

- ✦ **Real time Map reduce :**

Storm , Spark, Flink

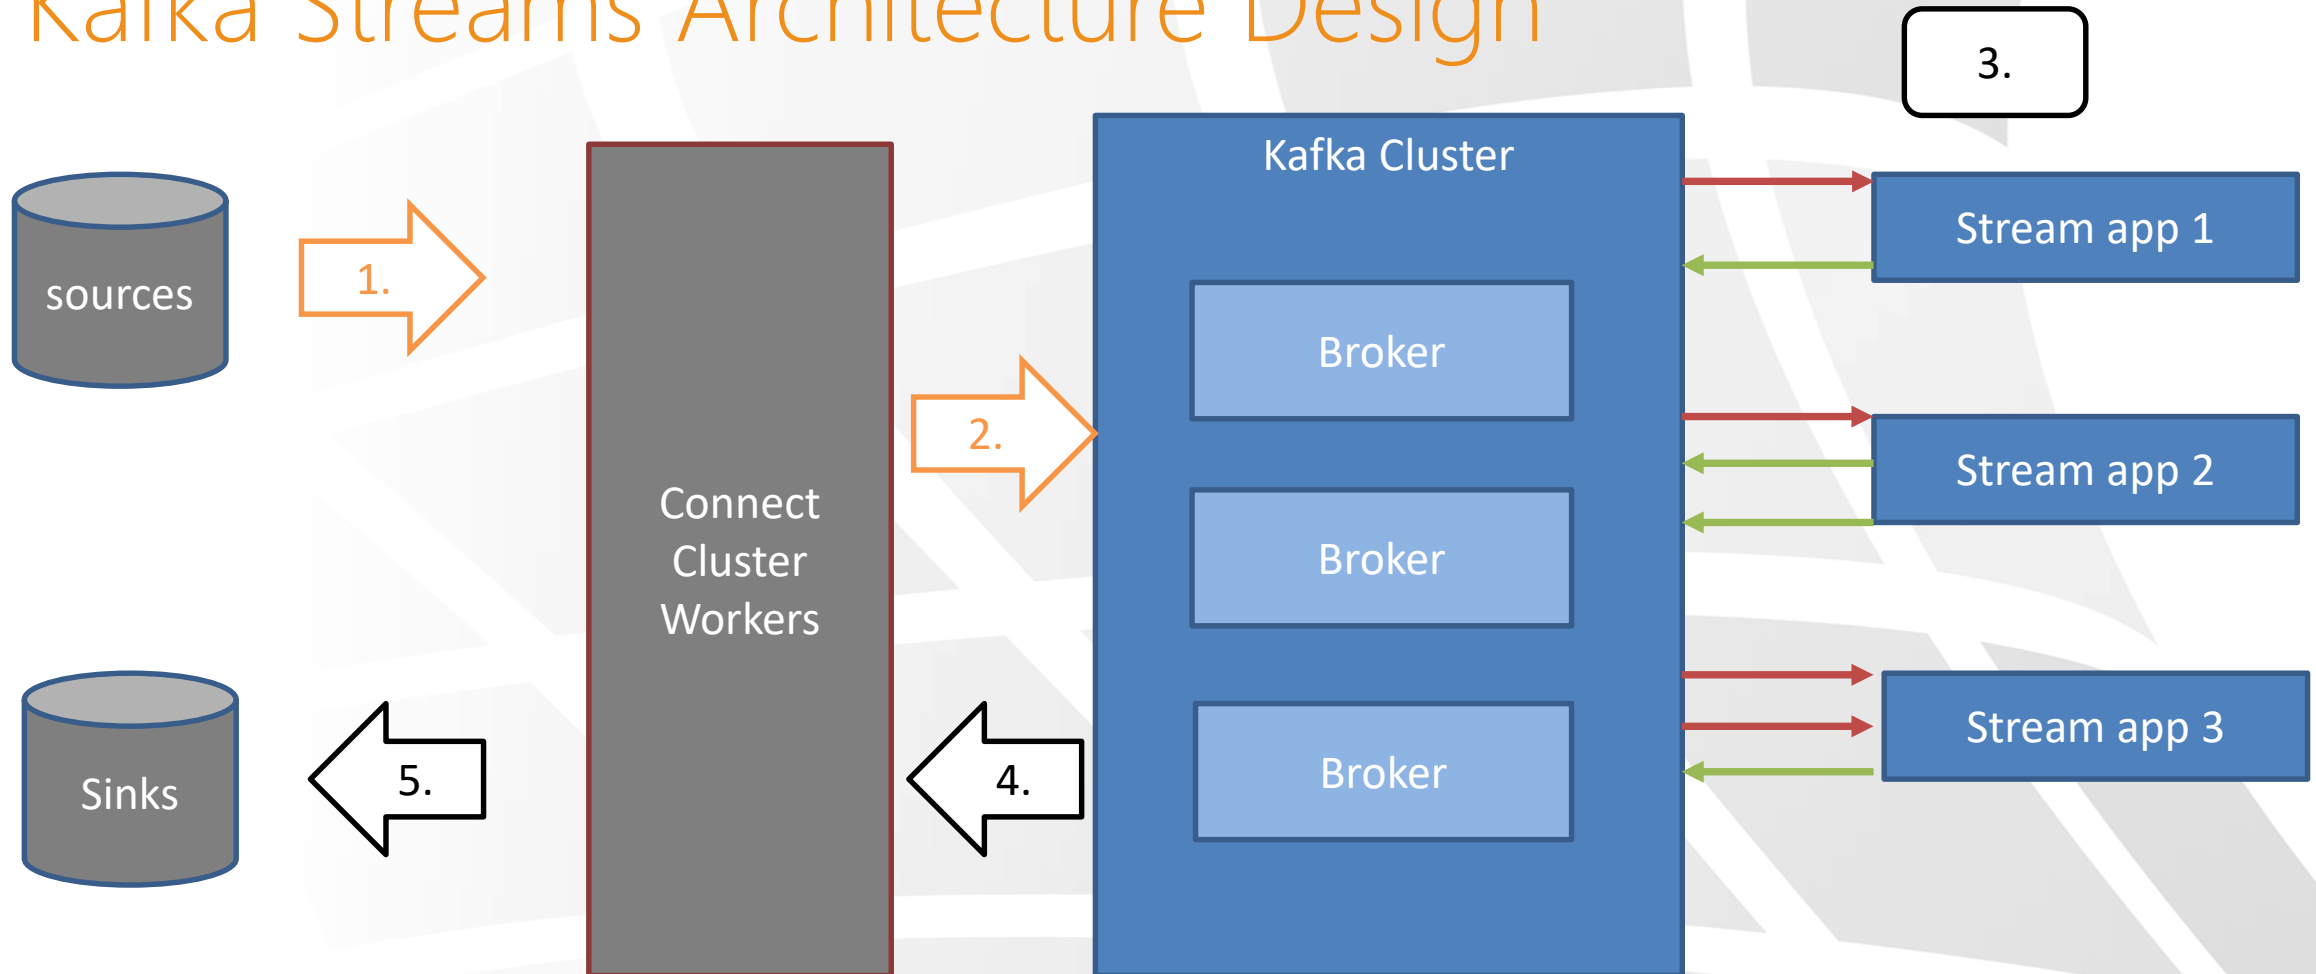
- ✦ They ran on dedicated clusters
    - ✦ Long running analytics, machine learning ,etc.

- ✦ **Even Driven micro-services :**

- ✦ The streaming platform is the central nervous system
    - ✦ Micro-services units acts as stream processing units
    - ✦ Input and output are always streams
    - ✦ Kafka Streams, Akka-Streams



# Kafka Streams Architecture Design




## Example usage

- ✦ We want to detect a fraud of credit cards which are detected by using the same Credit –card number **in a window of 5 minutes** within 2 different IP addresses
- ✦ We want to filter all tweets according to a key word or a topic and put the result back to kafka

# Stream Topology

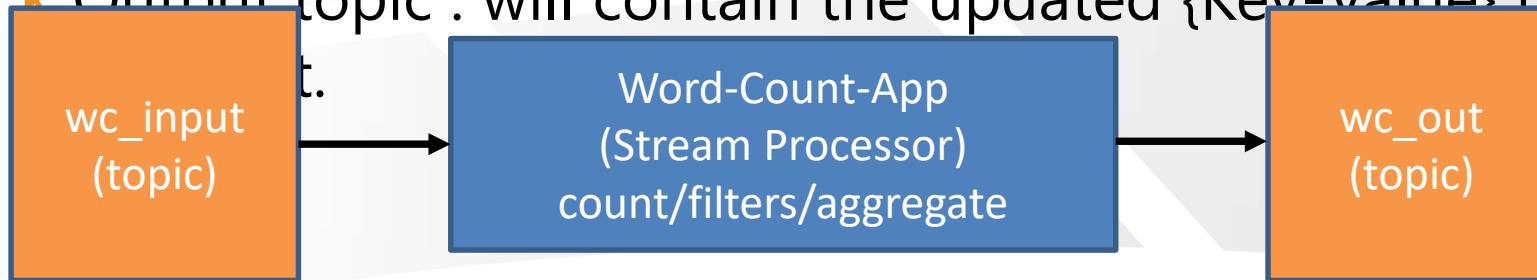
- ✦ A stream processor is the **node** in the topology, it represents the processing steps to transform data
- ✦ Arch represents the **stream**
- ✦ **Source Processor** – sends the data from one or more kafka topics.
- ✦ **Sink Processor** – does not have any more downstream processors.





# Building a word-count Pipeline

- ✦ We are going to build a simple pipeline of Kstreams that counts words
- ✦ Our java app will act as the Stream Processor and will interact with :
  - ✦ Input Topic : will contains sentences that will be break down to words
  - ✦ Output topic : will contain the updated {Key-Value} pair of each word and



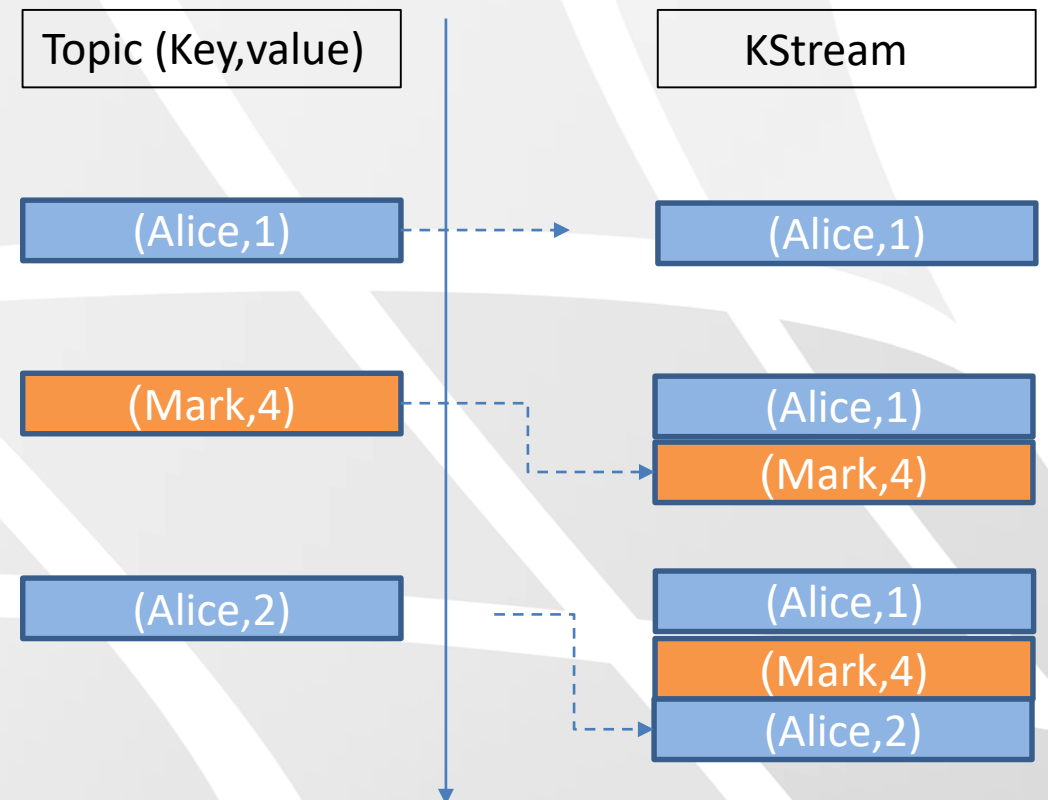
# Building a Word-Count Kafka Stream Pipeline

## Demo



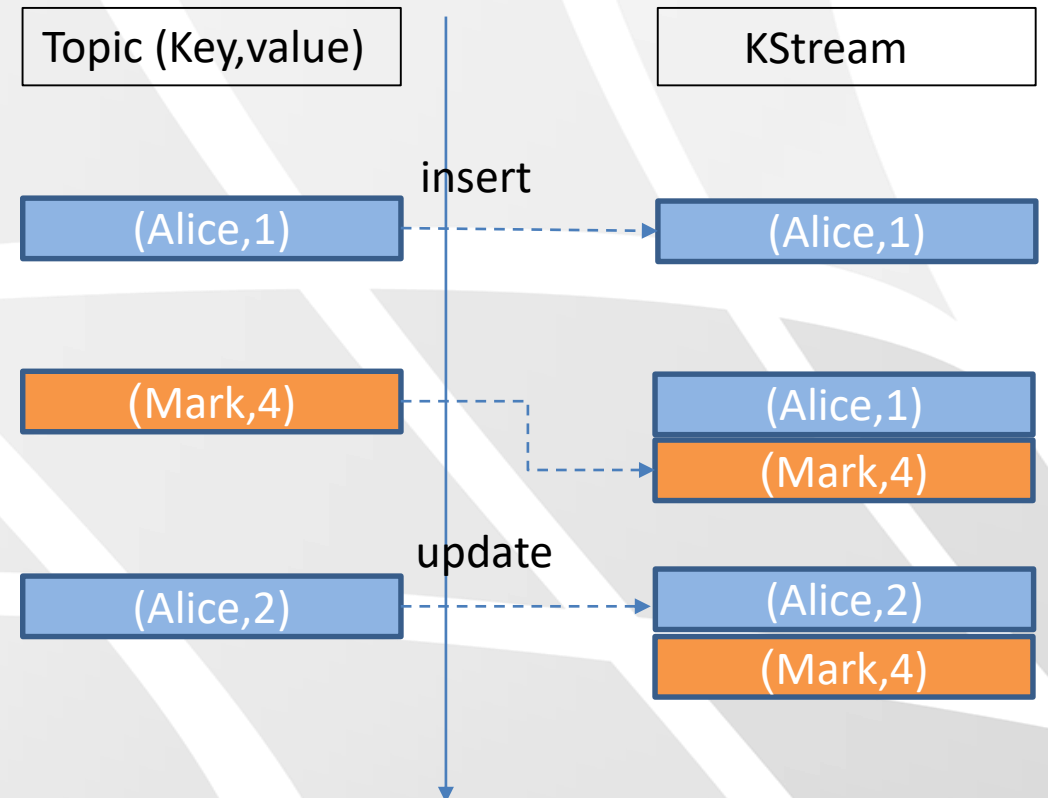
# Kstream

- ✦ An abstraction of data pipeline from Topics that supports :
  - ✦ All **inserts**
  - ✦ Similar to a log
  - ✦ Infinite
  - ✦ Unbounded data streams



# KTable

- ✦ All **Upserts** on non-null values
- ✦ Deletes on null value
- ✦ Similar to a table
- ✦ Parallel with log compact topics



# Demo Explained

KTable<String, Long>

KStream<String, Long>

all	1
-----	---

→ ("all", 1)

all	1
streams	1

→ ("streams", 1)

all	1
streams	1
lead	1

→ ("lead", 1)

all	1
streams	1
lead	1
to	1

→ ("to", 1)

all	1
streams	1
lead	1
to	1
kafka	1

→ ("kafka", 1)

KTable<String, Long>

KStream<String, Long>

all	1
streams	1
lead	1
to	1
kafka	1
hello	1

→ ("hello", 1)

all	1
streams	1
lead	1
to	1
kafka	2
hello	1

→ ("kafka", 2)

all	1
streams	2
lead	1
to	1
kafka	2
hello	1

Red numbers denote updates to existing table entries

→ ("streams", 2)

...

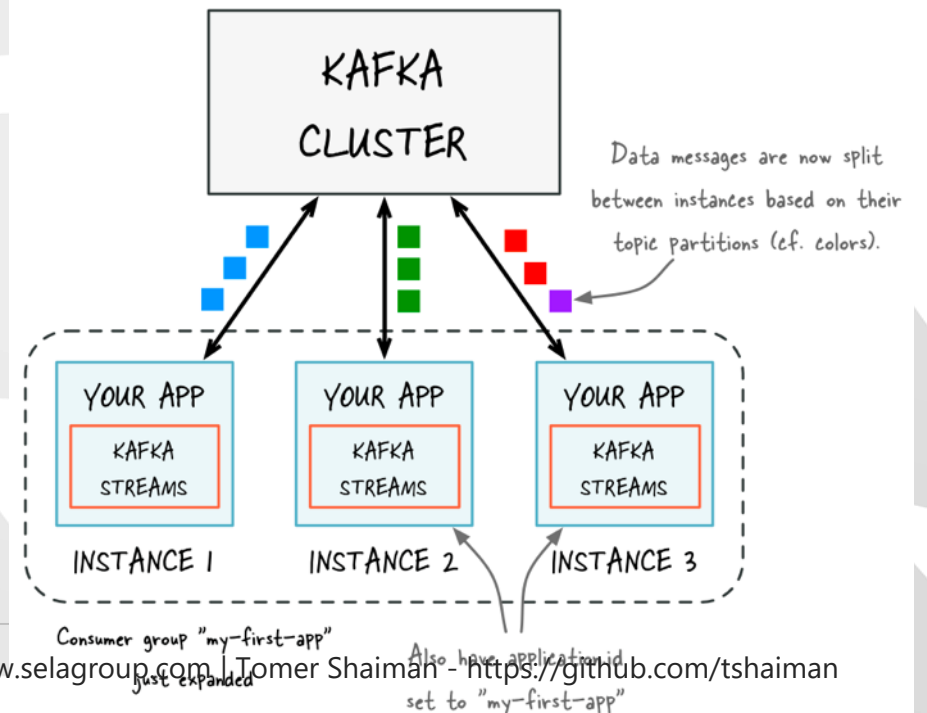
# Word Count Internal Topics

- ✦ Kafka Streams application will eventually create internal topics:
  - ✦ Repartition Topics – In case you transform your key in the stream, a repartition will happen
  - ✦ Changelog Topics- In case you create aggregation, Kafka streams will save compact data in these topics
- ✦ Internal Topics
  - ✦ are managed by Kafka Streams
  - ✦ are prefix by application.id parameter
  - ✦ **You should never delete, altered or publish to ,or change your app.id**

# Scaling our application

- ✦ Since our topic has **2 partitions**, we can launch up to 2 instances of the same app in parallel
- ✦ This is due to the fact that Kafka Streams application relies on **KafkaConsumer** which allows us to add more consumers to consumer group.

✦ Conclusion : **Scaling in Kafka Stream application does not require cluster**



# Running WordCount pipeline with filters

## Lab







## Module 07: Kafka Eco-System and Architecture



# Agenda

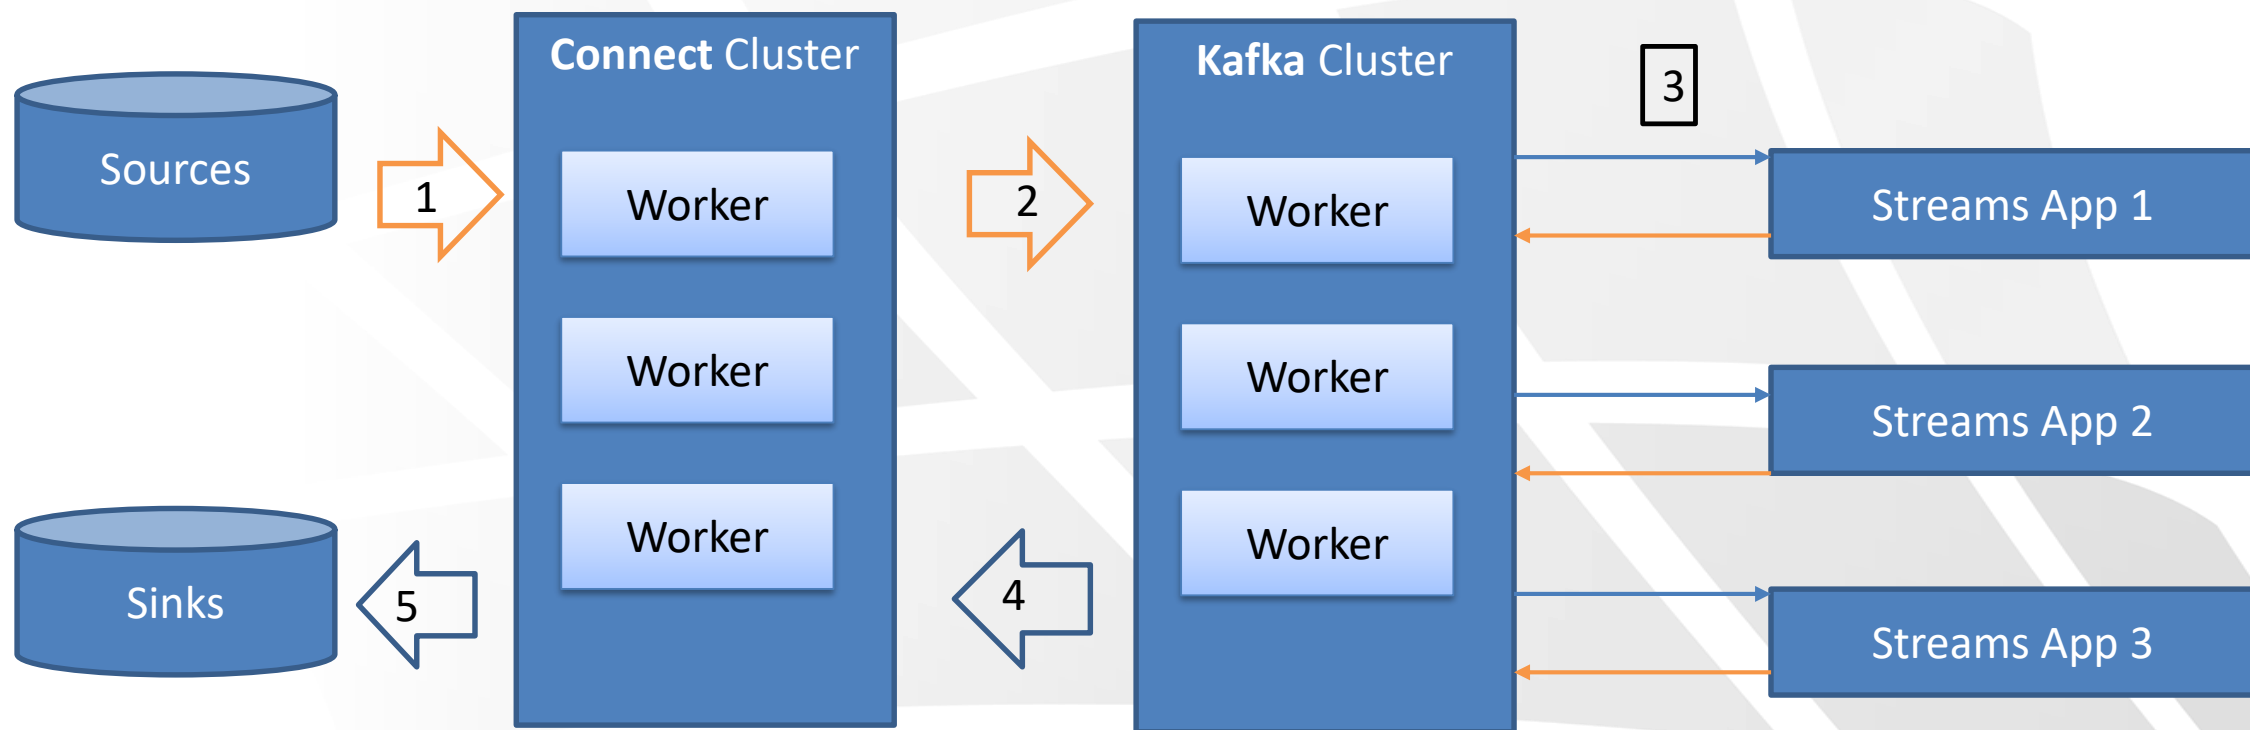
- ✦ Kafka Connect
- ✦ Kafka Schema Registry
- ✦ Architecture case Study
- ✦ Admins – Cluster Set-up
- ✦ Admins – Monitoring & Operations
- ✦ Start Kafka Differently

# Kafka Connect

- ✦ Simplify and improve getting data in and out of kafka.
- ✦ You're not the only one who read data from a DB/Stream and needs to send it to Kafka
- ✦ You're not the only one who needs to send data to Db/stream(Elasticsearch/ MongoDB/Sql/ PostgreSQL)
- ✦ Its hard to achieve fault -Tolerance, Idempotence, Distribution ordering



# Kafka Connect and Streams Architecture Design



## Scenario: Output Kafka Topic to a File (Kafka-Connect)

- ✦ It is common to output the content of a topic into a file
- ✦ Kafka Connect is a framework that provides scalable and reliable streaming of data to and from Apache Kafka.
- ✦ We will use a "File Sink" connector with configuration file as out "Sink"

```
#my-file-sink.properties config file
name=local-file-sink
connector.class=FileStreamSink
tasks.max=1
file=/tmp/my-file-sink.txt
topics=my-connect-test
```

# Kafka Connect –Cont. : Workers

✦ Processors that execute Kafka Connect connectors are called Workers

```
#bootstrap kafka servers
bootstrap.servers=localhost:9092

# specify input data format
key.converter=org.apache.kafka.connect.storage.StringConverter
value.converter=org.apache.kafka.connect.storage.StringConverter

# The internal converter used for offsets, most will always want to use the built-in default
internal.key.converter=org.apache.kafka.connect.json.JsonConverter
internal.value.converter=org.apache.kafka.connect.json.JsonConverter
internal.key.converter.schemas.enable=false
internal.value.converter.schemas.enable=false

# local file storing offsets and config data
offset.storage.file.filename=/tmp/connect.offsets
```

# Kafka-Connect – Running Workers

```
$ connect-standalone.sh my-standalone.properties my-file-sink.properties
```

- ✦ Launches the Kafka Connect Job On the sink properties file
- ✦ Worker manages its own offset into a file system without effecting and of the offset for other Consumer Groups
- ✦ Can be done for both ways : Producing a Data from File into Kafka Topic
- ✦ Extremely powerful for Porting /Testing / Automation tasks !

# Kafka Connect – Topic To File Sink

# Demo





# The need for Schema registry

- ✦ Kafka takes bytes as input and publishes them
- ✦ No data verification !
  - ✦ What if the producer produces bad data ?
  - ✦ What if a field gets rename ?
  - ✦ What if the format changes over time ?

## **The Consumer might break !**


- ✦ We need the data to be self – describable
- ✦ We need to evolve data without breaking downstream consumers
- ✦ We Need Schema + Schema registry

# The need for Schema Registry

- ✦ What if the kafka brokers were verifying the messages ?
- ✦ It would break what Kafka so good at :
  - ✦ Kafka doesn't parse or even read your data (no CPU usage)
  - ✦ Kafka takes bytes as input without loading them to memory
  - ✦ Kafka doesn't know if you wrote string, integer ,Boolean and does not parse the data.
- ✦ The Schema registry has to be separate component !
- ✦ Producers and Consumers needs to talk to it
- ✦ A common data format must be agreed upon

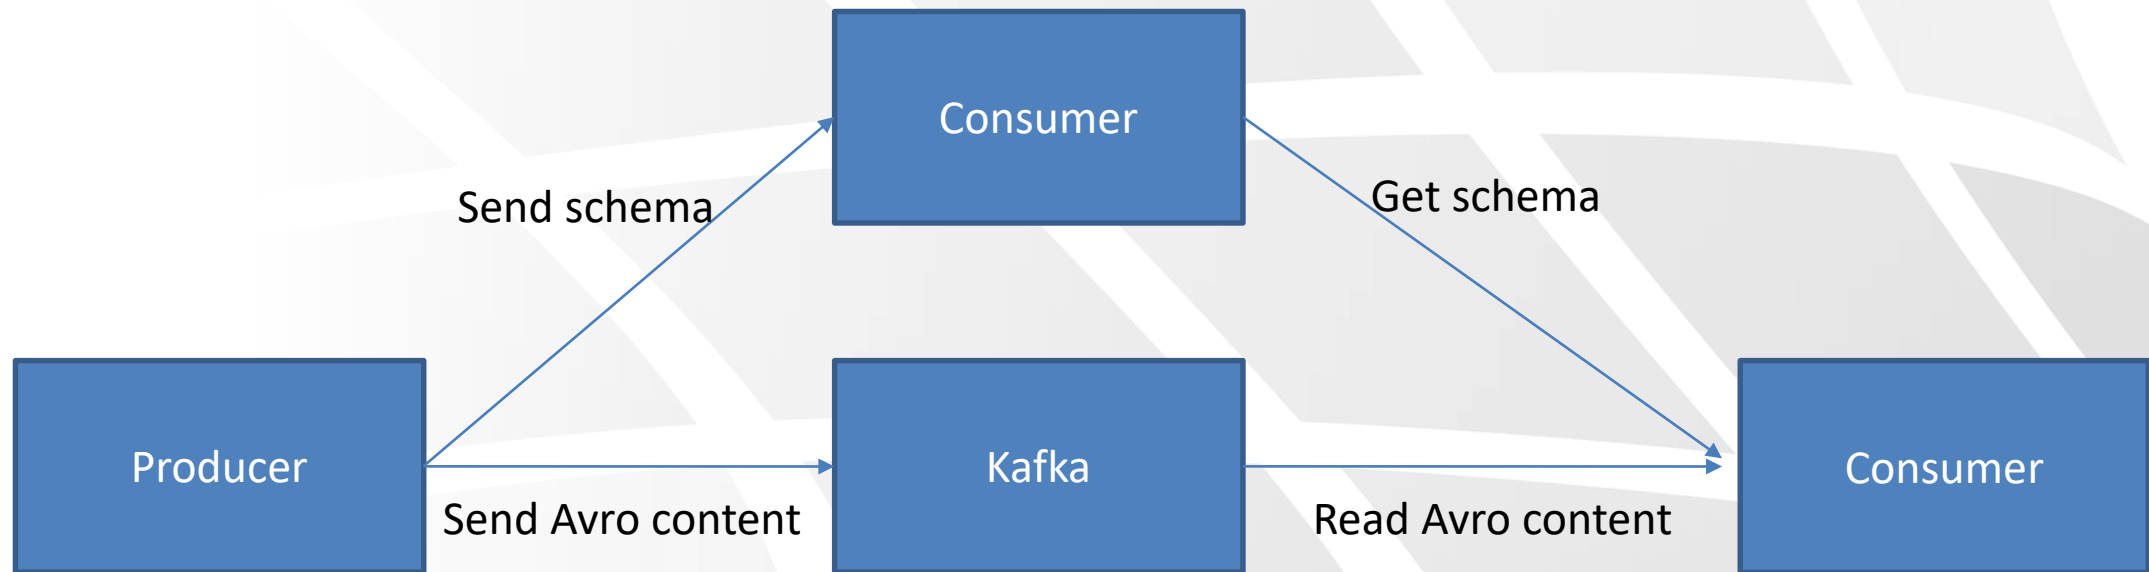


# Pipeline without Schema Registry




# Confluent Schema Registry Purpose

- ✦ Store and retrieve schemas for Producers / Consumers
- ✦ Enforce Backward / Forward /Full compatibility on topics
- ✦ Decrease the size of the payload of data sent to Kafka



# Schema Registry : gotchas

- ✦ There are many benefits for using Schema registry
- ✦ But it implies :
  - ✦ Set it up well
  - ✦ Make it highly available !
  - ✦ Partially change the producer and consumer code
- ✦ Apache **avro** is great but has a learning curve
- ✦ The Schema registry is free and open sourced ,created by Confluent

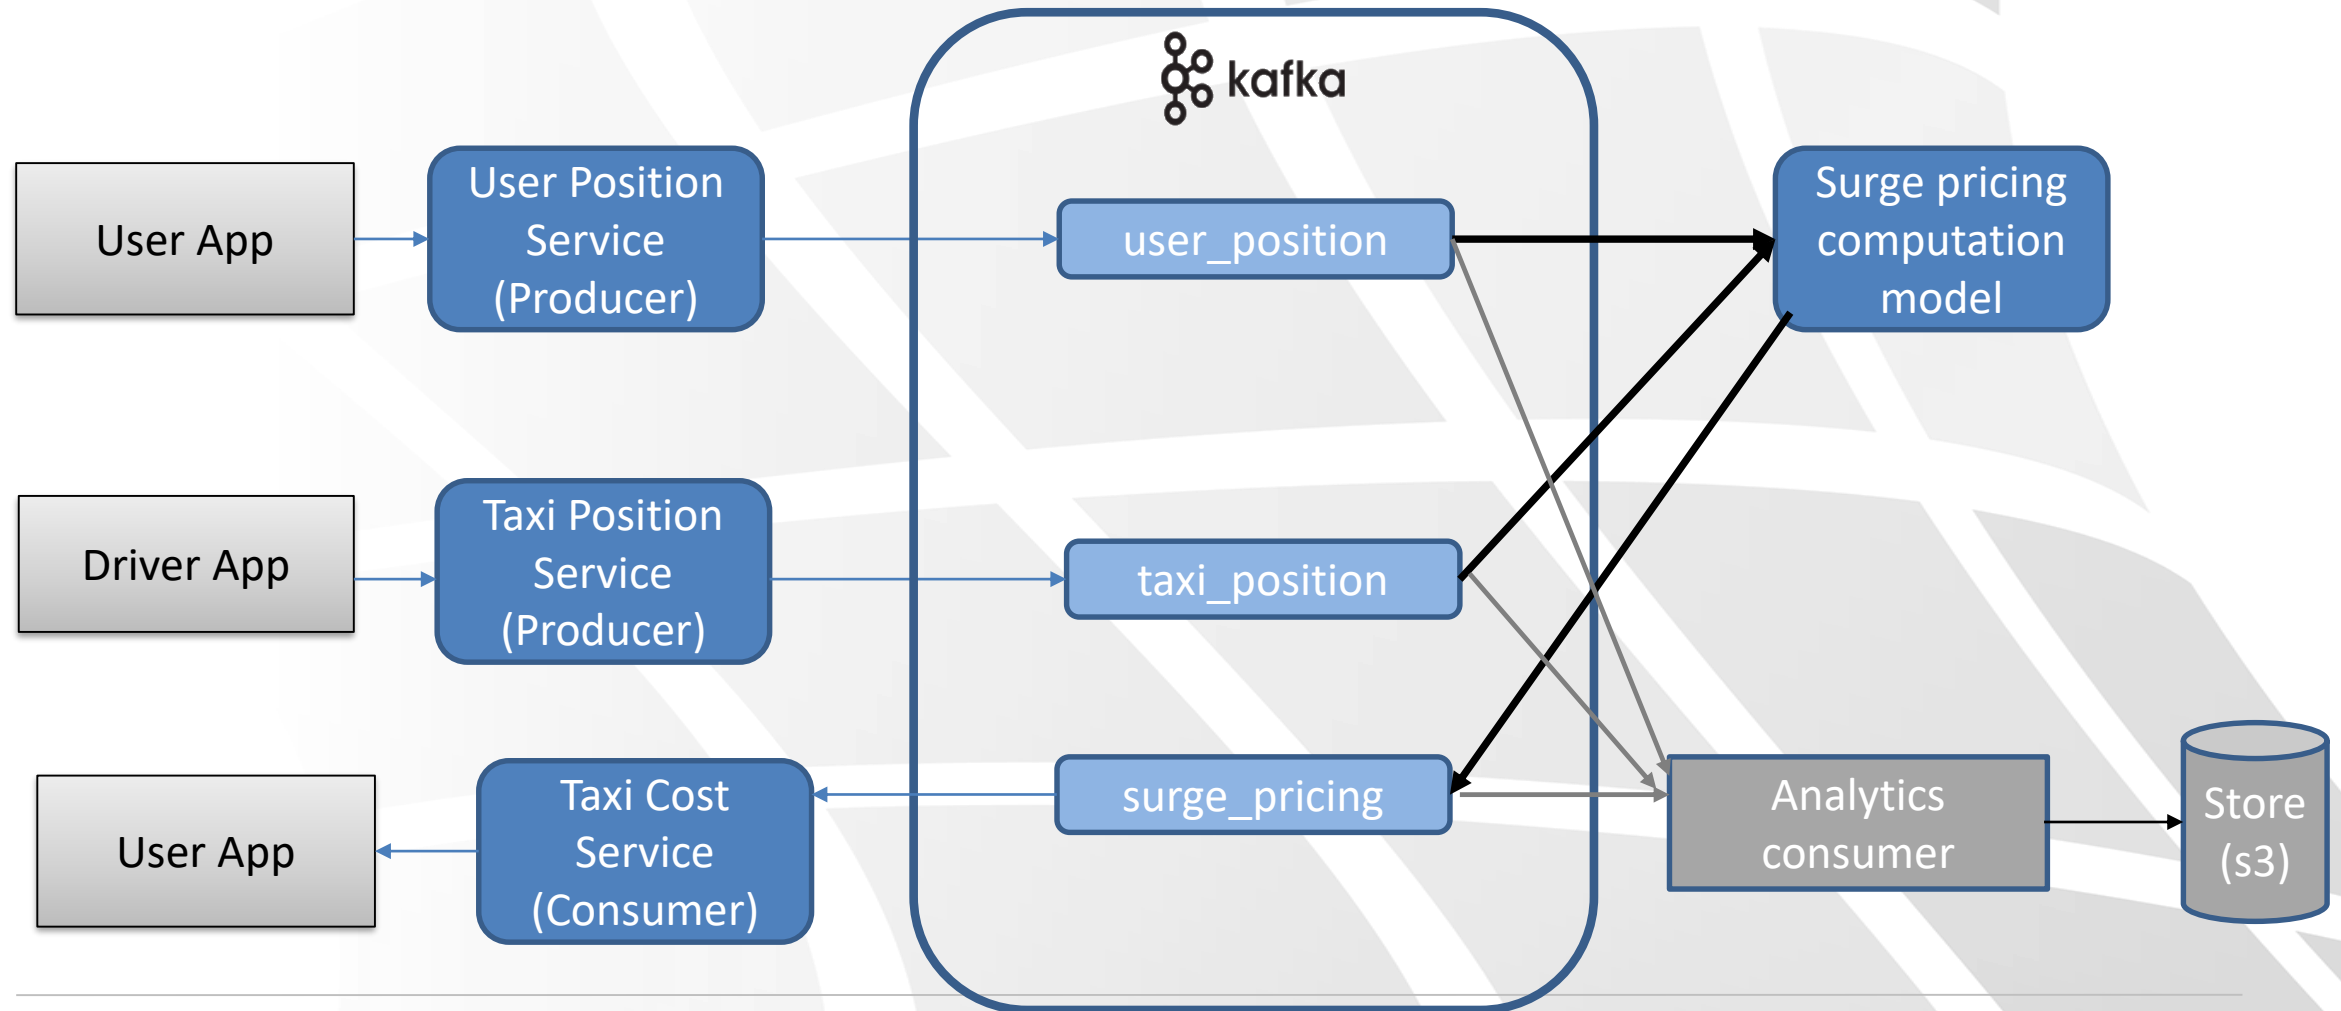


# Case Study – GetTaxi

- ✦ GetTaxi is a (fake) company that matches customers with taxi drivers on demand
- ✦ The Business wants the following:
  - ✦ Users should match with the closest driver
  - ✦ The pricing should “surge” if the numbers of drivers are low/high
  - ✦ All the position data before and during the ride should be stored in an analytics store so that the cost is computed correctly
- ✦ Take few minutes to come up with solution



# Case Study – GetTaxi



# Case Study – GetTaxi

## ✦ taxi\_position, user\_position topics :

- ✦ Can have multiple producers
- ✦ Should be highly distributed if high volume, topics > 30
- ✦ Key : could be "user\_id" / "taxi\_id"
- ✦ Data should not be stored for a long time

## ✦ surge\_pricing topic :

- ✦ The computation of it comes from Kafka streams app
- ✦ Surge pricing may be regional therefore the topic has high vo
- ✦ Maybe use "weather" or "events" for the Kstream App !

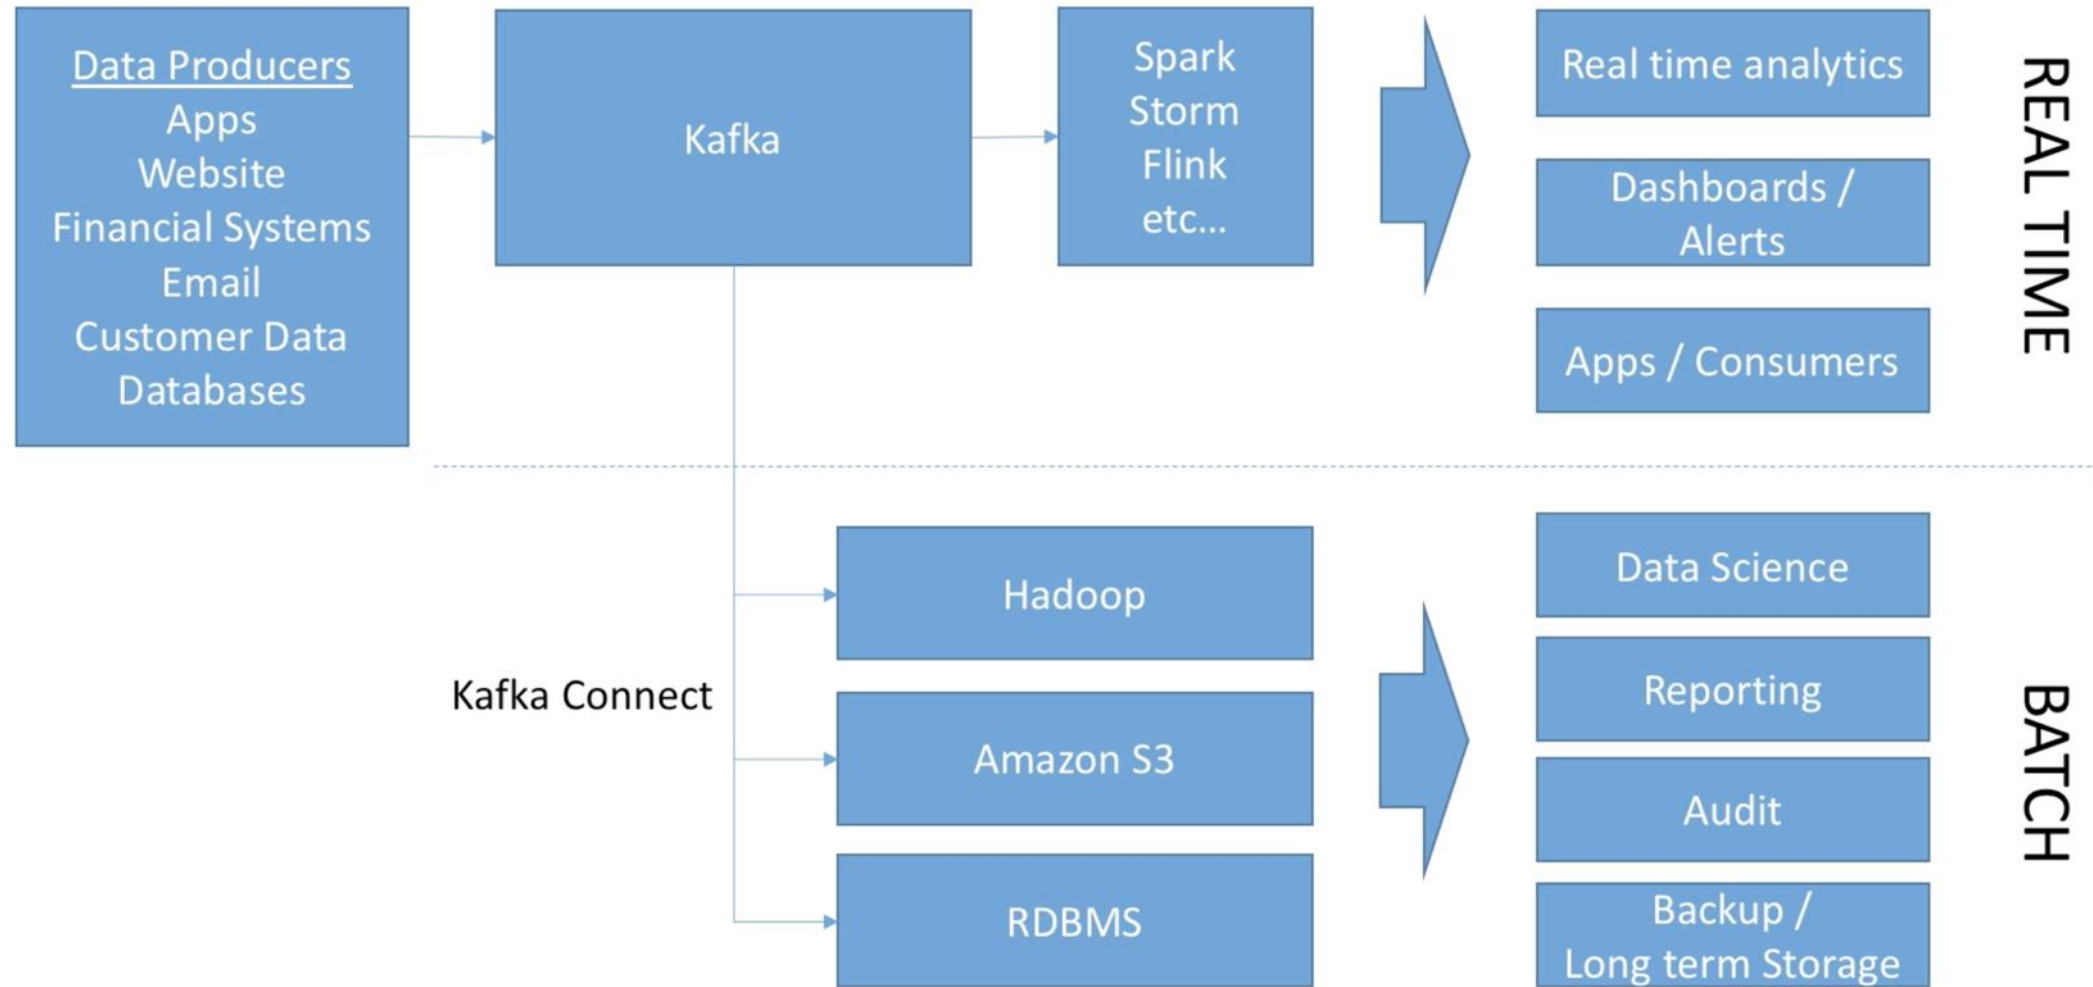




# Case Study – Big Data Ingestion

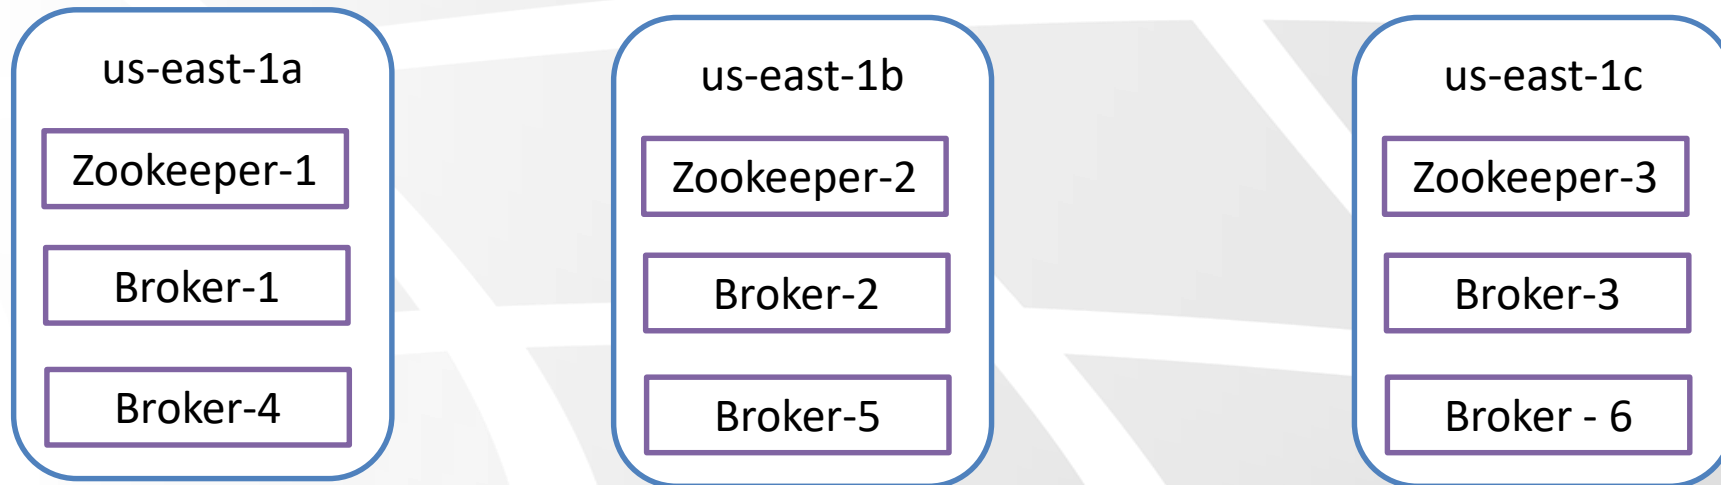
- ✦ Sometimes it is desired to have a “batch” pipeline for big-data analysis that occurs once in a while (machine learning/ analytics). This pipeline is considered “slow” ( S3, HDFS, ElasticSearch etc)
  - ✦ On the other hand , kafka can serve as a “near-real-time” /“fast” layer for processing large amount of data and react immediately (e.g : fraud detection, social media, brand-safety, ad-tech)
  - ✦ “Kafka-As-A-Front” to Big-Data ingestion is a common pattern, in such case kafka is like buffer before the store.
-

# Big Data Ingestion



# Kafka Cluster Setup – High Level Architecture

- ✦ You want multiple brokers in different data centers to distribute your load.
- ✦ You want a cluster of at least 3 zookeeper
- ✦ In AWS :



# Kafka Cluster

- ✦ Its not easy to setup an HA cluster
- ✦ You want to isolate the Zookeeper & Brokers on separate servers
- ✦ Monitoring needs to be implemented
- ✦ Operations have to be mastered
- ✦ You need a really good Kafka Admin !



- ✦ Alternative : Kafka-As-A-Service (on the web / Azure Kafka Service )

# Kafka Monitoring & Operations

- ✦ Kafka Exposes Metrics via JMX
- ✦ These metrics are highly important for monitoring kafka, and ensures the system behaves correctly under load
- ✦ Common Places to host the Kafka Metrics :
  - ✦ ELK (ElasticSearch + Kibana)
  - ✦ Datadog
  - ✦ NewRelic
  - ✦ Prometheus
  - ✦ Confluent Control Center
  - ✦ Influx Db

# Most Important Metrics

- ✦ Under Replicated Partitions : Number of Partitions that have problems with the ISR (In-Sync-Replica). May indicate a high load on the system
- ✦ Request Handlers : Utilization of threads for I/O, network,etc..overall utilization of Apache Kafka broker.
- ✦ Request timing : how long it takes to reply to requests. Lower is better
- ✦ Lag : the Lag offset of consumer group
- ✦ JVM Metrics : CPU / Memory /etc

# Monitoring Resources

✦ Excellent blog on how to configure InfluxDB to work with Kafka:

<https://softwaremill.com/monitoring-apache-kafka-with-influxdb-grafana/>

✦ Confluent Docs :

<https://docs.confluent.io/current/kafka/monitoring.html>


✦ Kafka Docs :

<https://kafka.apache.org/documentation/#monitoring>

# Operations

✦ Kafka Operations team must be able to perform the following tasks:

- ✦ Rolling Restart of brokers
- ✦ Updating Configurations
- ✦ Rebalancing partitions
- ✦ Increase replication factor
- ✦ Adding a broker
- ✦ Replace a broker
- ✦ Removing a broker
- ✦ Upgrading kafka cluster with zero downtime





# Start Kafka Differently

## ✦ Confluent CLI Tools :

<https://github.com/confluentinc/confluent-cli>

## ✦ Create a multi node cluster

- ✦ Clone the server.properties to be server1.properties
- ✦ Change ports
- ✦ Change log directory
- ✦ replicate data directory
- ✦ start 3 Instances of kafka-server-start.sh

## ✦ Using Docker-Compose command with the provided file:

✦ **\$ docker-compose up**



# Wrap Up



- ✦ Stream processing is Huge topic and we have covered many concepts
- ✦ You should try and build a simple pipeline and advance gradually
- ✦ Sys Admin and Security were not covered but are key factor of kafka cluster
- ✦ Think of How Kafka Connect can Assist you in your ETL process
- ✦ Kafka Streams can leverage your transformation and aggregation tasks from the Batch Processing into real-time streaming apps !
- ✦ Use Schema registry to protect your data and enable evolution and versioning !

# We Are here for you...



✦ Our Architects And Consultants has vast experience in :

- ✦ Demolish Monoliths
- ✦ Building Serverless Architecture on AWS
- ✦ PaaS and IaaS project on the cloud
- ✦ Reactive systems on open source technologies with variety of tech-stack and languages
- ✦ Streaming processing and Big Data Projects
- ✦ Docker Kubernetes and Service Mesh (Istio/OpenShift)



# Questions

